

# Overall

October 7, 2025

## 1 Multimodal Ensemble Architecture for Time Series Forecasting

- EAD
- Feature Engineering
- CNN-BiLSTM Modeling
- Transformer Modeling

### 1.1 Import Necessary Libraries

```
[1]: # Python Standard Libraries
import os
import csv
import math
import random
import unicodedata

# Data Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# NLP - NLTK
import nltk
nltk.download('vader_lexicon')
from nltk.sentiment.vader import SentimentIntensityAnalyzer

# Scikit-learn
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler, StandardScaler
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

# PyTorch
import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader

# TensorFlow / Keras
```

```

import tensorflow as tf
from tensorflow.keras.models import Model, Sequential
from tensorflow.keras.layers import (
    Input, Dense, Dropout, LSTM, Bidirectional,
    Conv1D, Conv2D, MaxPooling1D, MaxPooling2D,
    Flatten, GlobalAveragePooling1D, LayerNormalization,
    MultiHeadAttention, Add
)
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau, \
    ModelCheckpoint
import tensorflow.keras.backend as K
from tensorflow.keras.losses import Huber

# XGBoost
import xgboost as xgb
from xgboost import XGBRegressor

# Shap
import shap

```

```

[nltk_data] Downloading package vader_lexicon to
[nltk_data] /Users/yourth/nltk_data...
[nltk_data] Package vader_lexicon is already up-to-date!

```

## 1.2 Data Collections

```

[2]: raw_stocks = pd.read_csv('./data/raw/stock_yfinance.csv')
     raw_tweets = pd.read_csv('./data/raw/stock_tweets.csv')

```

## 1.3 Exploratory Data Analysis & Data Preprocessing

### 1.3.1 1. stock\_yfinance Dataset

```

[3]: raw_stocks.head()

```

```

[3]:
      Date      Open      High      Low      Close  Adj Close  \
0  2015-01-02  27.847500  27.860001  26.837500  27.332500  24.320429
1  2015-01-05  27.072500  27.162500  26.352501  26.562500  23.635286
2  2015-01-06  26.635000  26.857500  26.157499  26.565001  23.637505
3  2015-01-07  26.799999  27.049999  26.674999  26.937500  23.968958
4  2015-01-08  27.307501  28.037500  27.174999  27.972500  24.889906

      Volume Stock Name
0  212818400      AAPL
1  257142000      AAPL
2  263188400      AAPL
3  160423600      AAPL

```

4 237458000 AAPL

```
[4]: raw_stocks.tail()
```

```
[4]:
```

	Date	Open	High	Low	Close	Adj Close	\
6285	2019-12-24	27.890667	28.364668	27.512667	28.350000	28.350000	
6286	2019-12-26	28.527332	28.898666	28.423332	28.729334	28.729334	
6287	2019-12-27	29.000000	29.020666	28.407333	28.691999	28.691999	
6288	2019-12-30	28.586000	28.600000	27.284000	27.646667	27.646667	
6289	2019-12-31	27.000000	28.086000	26.805332	27.888666	27.888666	

	Volume	Stock Name
6285	120820500	TSLA
6286	159508500	TSLA
6287	149185500	TSLA
6288	188796000	TSLA
6289	154285500	TSLA

```
[ ]: raw_stocks.describe()
```

```
[ ]:
```

	Open	High	Low	Close	Adj Close	\
count	6290.000000	6290.000000	6290.000000	6290.000000	6290.000000	
mean	47.939449	48.373949	47.468337	47.943927	46.149473	
std	28.802247	28.991926	28.561164	28.793088	27.501038	
min	9.488000	10.331333	9.403333	9.578000	9.578000	
25%	26.413126	26.652000	26.131500	26.448125	24.851683	
50%	42.177500	42.528000	41.861500	42.233999	40.205023	
75%	58.738500	59.248500	58.206749	58.749249	57.247396	
max	159.449997	159.550003	158.220001	158.960007	151.738663	

	Volume
count	6.290000e+03
mean	7.857451e+07
std	6.388470e+07
min	7.425600e+06
25%	3.243302e+07
50%	6.051200e+07
75%	1.035403e+08
max	6.488252e+08

```
[6]: raw_stocks.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6290 entries, 0 to 6289
Data columns (total 8 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Date        6290 non-null  object
```

```

1  Open          6290 non-null  float64
2  High          6290 non-null  float64
3  Low           6290 non-null  float64
4  Close         6290 non-null  float64
5  Adj Close     6290 non-null  float64
6  Volume        6290 non-null  int64
7  Stock Name    6290 non-null  object
dtypes: float64(5), int64(1), object(2)
memory usage: 393.3+ KB

```

```
[7]: raw_stocks['Stock Name'].unique()
```

```
[7]: array(['AAPL', 'AMZN', 'GOOGL', 'MSFT', 'TSLA'], dtype=object)
```

```
[8]: stocks = raw_stocks.copy()

stocks['Date'] = pd.to_datetime(stocks['Date'])
stocks
```

```
[8]:
```

	Date	Open	High	Low	Close	Adj Close	\
0	2015-01-02	27.847500	27.860001	26.837500	27.332500	24.320429	
1	2015-01-05	27.072500	27.162500	26.352501	26.562500	23.635286	
2	2015-01-06	26.635000	26.857500	26.157499	26.565001	23.637505	
3	2015-01-07	26.799999	27.049999	26.674999	26.937500	23.968958	
4	2015-01-08	27.307501	28.037500	27.174999	27.972500	24.889906	
...	...	...	...	...	...	...	
6285	2019-12-24	27.890667	28.364668	27.512667	28.350000	28.350000	
6286	2019-12-26	28.527332	28.898666	28.423332	28.729334	28.729334	
6287	2019-12-27	29.000000	29.020666	28.407333	28.691999	28.691999	
6288	2019-12-30	28.586000	28.600000	27.284000	27.646667	27.646667	
6289	2019-12-31	27.000000	28.086000	26.805332	27.888666	27.888666	
	Volume	Stock Name					
0	212818400	AAPL					
1	257142000	AAPL					
2	263188400	AAPL					
3	160423600	AAPL					
4	237458000	AAPL					
...	...	...					
6285	120820500	TSLA					
6286	159508500	TSLA					
6287	149185500	TSLA					
6288	188796000	TSLA					
6289	154285500	TSLA					

```
[6290 rows x 8 columns]
```

### 1.3.2 2. stock\_tweets Dataset

```
[9]: raw_tweets.head()
```

```
[9]:
```

	Tweet ID	Writer	UTC \
0	550441509175443456	VisualStockRSRC	1420070457
1	550441672312512512	KeralaGuy77	1420070496
2	550441732014223360	DozenStocks	1420070510
3	550442977802207232	ShowDreamCar	1420070807
4	550443807834402816	i_Know_First	1420071005

	Tweet	Like Num	Stock Name \
0	1x21 made \$10,008 on \$AAPL -Check it out! htt...	1	AAPL
1	Insanity of today weirdo massive selling. \$aap...	0	AAPL
2	S&P100 #Stocks Performance \$HD \$LOW \$SBUX \$TGT...	0	AMZN
3	\$GM \$TSLA: Volkswagen Pushes 2014 Record Recal...	1	TSLA
4	Swing Trading: Up To 8.91% Return In 14 Days h...	1	AAPL

	Date
0	2015-01-01 00:00:57+00:00
1	2015-01-01 00:01:36+00:00
2	2015-01-01 00:01:50+00:00
3	2015-01-01 00:06:47+00:00
4	2015-01-01 00:10:05+00:00

```
[10]: raw_tweets.tail()
```

```
[10]:
```

	Tweet ID	Writer	UTC \
3943871	1212159838882533376	ShortingIsFun	1577836401
3943872	1212160015332728833	Commuternyc	1577836443
3943873	1212160410692046849	MoriaCrypto	1577836537
3943874	1212160410692046849	MoriaCrypto	1577836537
3943875	1212160477159206912	treabase	1577836553

	Tweet	Like Num \
3943871	In 2020 I may start Tweeting out positive news...	1
3943872	Patiently Waiting for the no twitter sitter tw...	5
3943873	I don't discriminate. I own both \$aapl and \$ms...	1
3943874	I don't discriminate. I own both \$aapl and \$ms...	1
3943875	\$AAPL #patent 10,522,475 Vertical interconnect...	0

	Stock Name	Date
3943871	TSLA	2019-12-31 23:53:21+00:00
3943872	TSLA	2019-12-31 23:54:03+00:00
3943873	MSFT	2019-12-31 23:55:37+00:00
3943874	AAPL	2019-12-31 23:55:37+00:00
3943875	AAPL	2019-12-31 23:55:53+00:00

```
[11]: raw_tweets.info(show_counts=True)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3943876 entries, 0 to 3943875
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Tweet ID    3943876 non-null  int64
1   Writer      3895635 non-null  object
2   UTC         3943876 non-null  int64
3   Tweet       3943876 non-null  object
4   Like Num    3943876 non-null  int64
5   Stock Name  3943876 non-null  object
6   Date        3943876 non-null  object
dtypes: int64(3), object(4)
memory usage: 210.6+ MB
```

```
[12]: tweets = raw_tweets.copy()

tweets = tweets.drop(columns=['Tweet ID', 'Writer', 'UTC', 'Like Num'])
tweets = tweets.dropna(subset=['Date', 'Tweet'])
tweets.isna().sum()
```

```
[12]: Tweet      0
      Stock Name  0
      Date       0
      dtype: int64
```

```
[13]: # Make sure Date is in datetime format
tweets['Date'] = pd.to_datetime(tweets['Date'], errors='coerce')
tweets['Date'] = tweets['Date'].dt.date

tweets.head()
```

```
[13]:
```

	Tweet	Stock Name	Date
0	1x21 made \$10,008 on \$AAPL -Check it out! htt...	AAPL	2015-01-01
1	Insanity of today weirdo massive selling. \$aap...	AAPL	2015-01-01
2	S&P100 #Stocks Performance \$HD \$LOW \$SBUX \$TGT...	AMZN	2015-01-01
3	\$GM \$TSLA: Volkswagen Pushes 2014 Record Recal...	TSLA	2015-01-01
4	Swing Trading: Up To 8.91% Return In 14 Days h...	AAPL	2015-01-01

```
[14]: tweets.info(show_counts=True)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3943876 entries, 0 to 3943875
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Tweet       3943876 non-null  object
```

```

1   Stock Name  3943876 non-null  object
2   Date       3943876 non-null  object
dtypes: object(3)
memory usage: 90.3+ MB

```

```

[15]: tweet_stock_count = raw_tweets['Stock Name'].value_counts()

perc = tweet_stock_count / tweet_stock_count.sum()
dstr = pd.DataFrame({'Count': tweet_stock_count, 'Percentage': perc}).
    ↪reset_index()
dstr = dstr.rename(columns={'index': 'Stock Name'})
dstr.head()

```

```

[15]:   Stock Name    Count  Percentage
0      AAPL  1425013    0.361323
1      TSLA  1096868    0.278119
2      AMZN   718715    0.182236
3      MSFT   375711    0.095264
4      GOOGL   327569    0.083058

```

```

[16]: # Set a threshold percentage
threshold = 0.02 # 2%
colors = ['#ff9999', '#66b3ff', '#99ff99', '#ffcc99', '#c2c2f0', '#ffb3e6',
    ↪'#c4e17f', '#76d7c4', '#f7c6c7']


# Create a new DataFrame grouping smaller slices
dstr['Grouped'] = dstr.apply(lambda row: 'Other' if row['Percentage'] <
    ↪threshold else row['Stock Name'], axis=1)

dstr_grouped = dstr.groupby(
    dstr['Grouped'].where(dstr['Grouped'] != 'Other', 'Other')
).agg({'Count': 'sum'}).reset_index()

# Plot
plt.figure(figsize=(8, 8))
plt.pie(
    dstr_grouped['Count'],
    labels=dstr_grouped['Grouped'],
    autopct='%1.1f%%',
    startangle=140,
    colors=colors,
    textprops={'fontsize': 12}
)
plt.title('Stock Tweets Distribution', fontsize=14, pad=30)
plt.axis('equal')
plt.show()

```

Stock Tweets Distribution



```
[17]: stocks.head()
```

```
[17]:
```

	Date	Open	High	Low	Close	Adj Close	\
0	2015-01-02	27.847500	27.860001	26.837500	27.332500	24.320429	
1	2015-01-05	27.072500	27.162500	26.352501	26.562500	23.635286	
2	2015-01-06	26.635000	26.857500	26.157499	26.565001	23.637505	
3	2015-01-07	26.799999	27.049999	26.674999	26.937500	23.968958	
4	2015-01-08	27.307501	28.037500	27.174999	27.972500	24.889906	

	Volume	Stock Name
0	212818400	AAPL
1	257142000	AAPL
2	263188400	AAPL



```
3 160423600      AAPL
4 237458000      AAPL
```

```
[18]: tweets.head()
```

```
[18]:
```

	Tweet	Stock Name	Date
0	1x21 made \$10,008 on \$AAPL -Check it out! htt...	AAPL	2015-01-01
1	Insanity of today weirdo massive selling. \$aap...	AAPL	2015-01-01
2	S&P100 #Stocks Performance \$HD \$LOW \$SBUX \$TGT...	AMZN	2015-01-01
3	\$GM \$TSLA: Volkswagen Pushes 2014 Record Recal...	TSLA	2015-01-01
4	Swing Trading: Up To 8.91% Return In 14 Days h...	AAPL	2015-01-01

```
[19]: def find_last_trading_day(x, trading_days):
    """
    x trading_days x

    NaT
    """
    eligible_days = trading_days[trading_days <= x]
    if not eligible_days.empty:
        return eligible_days.max()
    else:
        return pd.NaT

def map_to_last_trading_day(tweet_dates, trading_days):
    """
    tweet_dates: Series DatetimeIndex
    trading_days:

    Series
    """
    trading_days = pd.to_datetime(sorted(set(trading_days)))

    # apply + lambda
    mapped_dates = tweet_dates.apply(lambda x: find_last_trading_day(x,
↪trading_days))

    # NaT
    valid_mask = mapped_dates.notna()

    return mapped_dates[valid_mask]
```

```
[20]: #
tweets['Date'] = pd.to_datetime(tweets['Date'])
stocks['Date'] = pd.to_datetime(stocks['Date'])

#
trading_days = stocks['Date'].unique()

#
mapped_dates = map_to_last_trading_day(tweets['Date'], trading_days)

#     tweets
tweets = tweets.loc[mapped_dates.index].copy()
tweets['Trading Date'] = mapped_dates
```

```
[21]: tweets['Date'] = tweets['Trading Date']
tweets = tweets.drop(columns=['Trading Date']) #

tweets.head()
```

```
[21]:
```

		Tweet	Stock Name	Date
628	S&P100 #Stocks Performance \$HD \$LOW \$SBUX \$TGT...		AMZN	2015-01-02
629	Will Audi's Electric Q7 Cause \$TSLA Model X Ba...		TSLA	2015-01-02
630	Either way you're a winnah. RT @dbbrakebill: i...		AAPL	2015-01-02
631	@Weeklyoptions http://Weeklyoptionplays.com we...		AAPL	2015-01-02
632	Cash flow machine. RT @themandotcom: \$MSFT, wh...		MSFT	2015-01-02

## 1.4 Feature Construction

### 1.4.1 1. Sentiment Analysis

```
[22]: def vader_sentiment_scores(df, text_col='Tweet'):

    sentiment_analyzer = SentimentIntensityAnalyzer()
    df = df.copy()

    tweets['Vader_Negative'] = np.nan
    tweets['Vader_Neutral'] = np.nan
    tweets['Vader_Positive'] = np.nan
    tweets['Vader_Polarity'] = np.nan

    for indx, row in df.iterrows():
        try:
            # Normalize the text to ASCII
            text = unicodedata.normalize('NFKD', row[text_col])
            sentiment = sentiment_analyzer.polarity_scores(text)

            df.at[indx, 'Vader_Negative'] = sentiment['neg']
            df.at[indx, 'Vader_Neutral'] = sentiment['neu']
```

```

df.at[indx, 'Vader_Positive'] = sentiment['pos']
df.at[indx, 'Vader_Polarity'] = sentiment['compound']

except TypeError:
    print(f"TypeError on row {indx}: {row[text_col]}")
    break

return df

```

```
[23]: tweets_sent = vader_sentiment_scores(tweets)
```

```
[24]: tweets_sent.head()
```

```
[24]:
```

	Tweet	Stock Name	Date \
628	S&P100 #Stocks Performance \$HD \$LOW \$SBUX \$TGT...	AMZN	2015-01-02
629	Will Audi's Electric Q7 Cause \$TSLA Model X Ba...	TSLA	2015-01-02
630	Either way you're a winnah. RT @dbbrakebill: i...	AAPL	2015-01-02
631	@Weeklyoptions http://Weeklyoptionplays.com we...	AAPL	2015-01-02
632	Cash flow machine. RT @themandotcom: \$MSFT, wh...	MSFT	2015-01-02

	Vader_Negative	Vader_Neutral	Vader_Positive	Vader_Polarity
628	0.000	1.000	0.000	0.0000
629	0.000	1.000	0.000	0.0000
630	0.155	0.845	0.000	-0.3736
631	0.000	0.756	0.244	0.6369
632	0.000	1.000	0.000	0.0000

```
[25]: def plot_kde(df, polarity_col, threshold=0, stock_col='Stock Name'):

    plt.figure(figsize=(16, 8))

    sns.kdeplot(data=df, x=polarity_col, hue=stock_col, fill=True)


    if threshold != 0:
        # Vertical threshold lines
        plt.axvline(-threshold, color='red', linestyle='--', label=f'Lower_
↳Threshold -{threshold}')
        plt.axvline(threshold, color='blue', linestyle='--', label=f'Upper_
↳Threshold +{threshold}')

    # Final touches
    plt.title(f"{polarity_col} Distribution with ±{threshold} Thresholds",
↳fontsize=16)
    plt.xlabel(polarity_col, fontsize=14)
    plt.ylabel("Density", fontsize=14)
    plt.grid(True)
    plt.tight_layout()

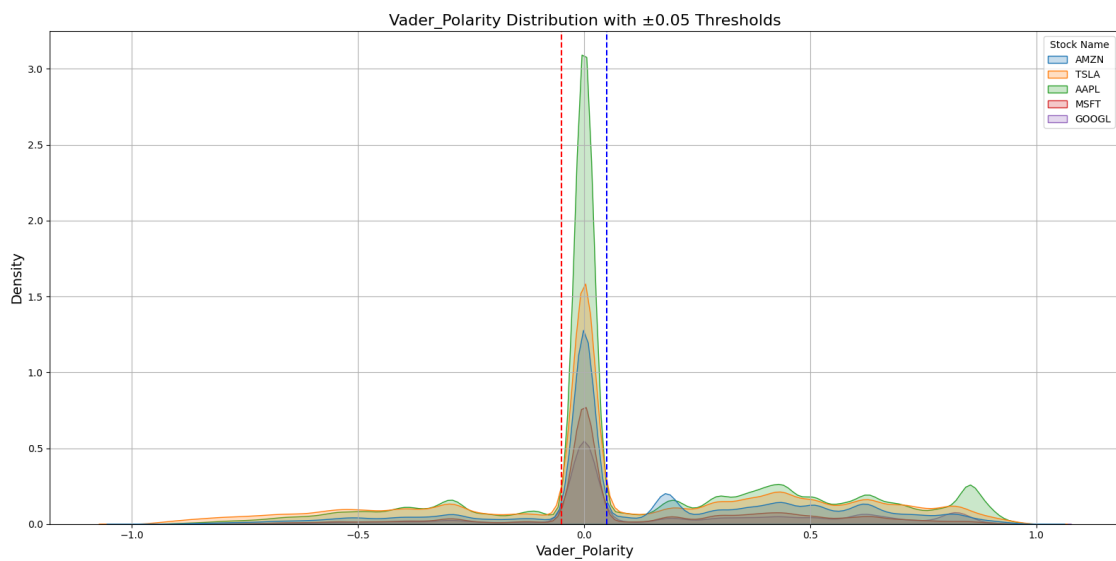
```

```
plt.show()
```

```
[26]: plot_kde(tweets_sent, polarity_col='Vader_Polarity')
```



```
[27]: threshold = 0.05  
plot_kde(tweets_sent, polarity_col='Vader_Polarity', threshold=threshold)
```



```
[28]: def polarity_filter_by_threshold(df, threshold, polarity_col='Vader_Polarity',  
    ↪ date_col='Date', verbose=True):
```

```

df = df.copy()
df[date_col] = pd.to_datetime(df[date_col])

# Apply polarity threshold filter
filtered_df = df[
    (df[polarity_col] >= threshold) |
    (df[polarity_col] <= -threshold)
]

# Compute stats
total_dates = df[date_col].nunique()
remaining_dates = filtered_df[date_col].nunique()
dates_lost = total_dates - remaining_dates

if verbose:
    print(f"-> Total unique dates before filtering: {total_dates}")
    print(f"+> Remaining unique dates after filtering: {remaining_dates}")
    print(f">> Dates lost: {dates_lost}")

return filtered_df

```

```
[29]: tweets_filtered = polarity_filter_by_threshold(tweets_sent, threshold)
```

```

-> Total unique dates before filtering: 1258
+> Remaining unique dates after filtering: 1258
>> Dates lost: 0

```

```

[30]: def plot_daily_sentiment(df, company, sentiment_col, date_col='Date',
    ↪ company_col='Stock Name'):

    # Convert date column to datetime if needed
    df = df.copy()
    df[date_col] = pd.to_datetime(df[date_col])

    # Group and compute daily average sentiment
    daily_sentiment = (
        df.groupby([date_col, company_col])[sentiment_col]
        .mean()
        .reset_index(name='Avg_Sentiment')
    )

    # Filter for selected company
    company_df = daily_sentiment[daily_sentiment[company_col] == company]

    # Plot
    plt.figure(figsize=(22, 10))

```

```

plt.plot(company_df[date_col], company_df['Avg_Sentiment'], color='blue',
↪marker='o')
plt.title(f"{company} - Daily Average {sentiment_col}", fontsize=16)
plt.xlabel("Date", fontsize=14)
plt.ylabel("Avg_Sentiment", fontsize=14)
plt.axhline(0, color='gray', linestyle='--')
plt.grid(True)
plt.tight_layout()
plt.xticks(rotation=45)
plt.show()

```

```
[31]: tweets_sent.head()
```


```
[31]:
```

	Tweet	Stock Name	Date \
628	S&P100 #Stocks Performance \$HD \$LOW \$SBUX \$TGT...	AMZN	2015-01-02
629	Will Audi's Electric Q7 Cause \$TSLA Model X Ba...	TSLA	2015-01-02
630	Either way you're a winnah. RT @dbbrakebill: i...	AAPL	2015-01-02
631	@Weeklyoptions http://Weeklyoptionplays.com we...	AAPL	2015-01-02
632	Cash flow machine. RT @themandotcom: \$MSFT, wh...	MSFT	2015-01-02


  

	Vader_Negative	Vader_Neutral	Vader_Positive	Vader_Polarity
628	0.000	1.000	0.000	0.0000
629	0.000	1.000	0.000	0.0000
630	0.155	0.845	0.000	-0.3736
631	0.000	0.756	0.244	0.6369
632	0.000	1.000	0.000	0.0000

```
[32]: plot_daily_sentiment(tweets_sent, company='TSLA',
↪sentiment_col='Vader_Polarity')
```



```
[33]: plot_daily_sentiment(tweets_filtered, company='TSLA',
↪sentiment_col='Vader_Polarity')
```



```
[34]: tweets_filtered.head()
```

```
[34]:
```

		Tweet	Stock Name	Date	\
630		Either way you're a winnah. RT @dbbrakebill: i...	AAPL	2015-01-02	
631		@Weeklyoptions http://Weeklyoptionplays.com we...	AAPL	2015-01-02	
633		.@BenBajarin @counternotions @GlennF Thankful ...	AAPL	2015-01-02	
634		perfectly trading the S&P 500 in 2014 \$FB \$MU ...	AMZN	2015-01-02	
637		@KyleRohde @thehilker It could be, but that's ...	AMZN	2015-01-02	

	Vader_Negative	Vader_Neutral	Vader_Positive	Vader_Polarity
630	0.155	0.845	0.000	-0.3736
631	0.000	0.756	0.244	0.6369
633	0.000	0.802	0.198	0.5719
634	0.000	0.781	0.219	0.6369
637	0.103	0.641	0.255	0.7351

```
[35]: tweets_filtered.to_csv("./data/filtered/tweets_filtered.csv", index=False)
```

## 1.4.2 2. Calculate Technical Indicators

```
[36]: # RSI Calculation
def calculate_rsi(series, window=14):
    delta = series.diff()
    gain = delta.where(delta > 0, 0).rolling(window=window).mean()
    loss = -delta.where(delta < 0, 0).rolling(window=window).mean()
    rs = gain / loss
    rsi = 100 - (100 / (1 + rs))
```

```
return rsi
```

```
[37]: def get_technical_indicators(df):
    df = df.copy()
    # Trend
    df['SMA_5'] = df['Close'].rolling(window=5).mean()
    df['SMA_20'] = df['Close'].rolling(window=20).mean()

    # Bollinger Bands
    df['BB_Mid'] = df['SMA_5']
    df['BB_Std'] = df['Close'].rolling(window=20).std()
    df['BB_Upper'] = df['BB_Mid'] + 2 * df['BB_Std']
    df['BB_Lower'] = df['BB_Mid'] - 2 * df['BB_Std']

    # RSI
    df['RSI_14'] = calculate_rsi(df['Close'], window=14)

    # Log Return
    df['Log_Return'] = np.log(df['Close'] / df['Close'].shift(1))

    # OBV
    df['OBV'] = (np.sign(df['Close'].diff()) * df['Volume']).fillna(0).cumsum()

    # Lag Features
    df['Prev_Close'] = df['Close'].shift(1)
    df['Prev_Volume'] = df['Volume'].shift(1)

    # Time Features
    df['DayOfWeek'] = df.index.dayofweek
    df['Month'] = df.index.month

    # Clean
    df = df.bfill().ffill()

    return df
```

```
[38]: stocks.head()
```

```
[38]:
```

	Date	Open	High	Low	Close	Adj Close	\
0	2015-01-02	27.847500	27.860001	26.837500	27.332500	24.320429	
1	2015-01-05	27.072500	27.162500	26.352501	26.562500	23.635286	
2	2015-01-06	26.635000	26.857500	26.157499	26.565001	23.637505	
3	2015-01-07	26.799999	27.049999	26.674999	26.937500	23.968958	
4	2015-01-08	27.307501	28.037500	27.174999	27.972500	24.889906	

	Volume	Stock Name
0	212818400	AAPL



```

1 257142000 AAPL
2 263188400 AAPL
3 160423600 AAPL
4 237458000 AAPL

```

```

[39]: stocks = stocks.set_index('Date').sort_index()
stocks = get_technical_indicators(stocks)

stocks.head()

```

```

[39]:
      Open      High      Low      Close  Adj Close      Volume \
Date
2015-01-02  27.847500  27.860001  26.837500  27.332500  24.320429  212818400
2015-01-02  14.858000  14.883333  14.217333  14.620667  14.620667   71466000
2015-01-02  26.629999  26.790001  26.393999  26.477501  26.351515  26480000
2015-01-02  46.660000  47.419998  46.540001  46.759998  40.072136  27913900
2015-01-02  15.629000  15.737500  15.348000  15.426000  15.426000  55664000

```

```

      Stock Name      SMA_5      SMA_20      BB_Mid      BB_Std      BB_Upper \
Date
2015-01-02      AAPL  26.123333  25.620192  26.123333  11.873179  49.228223
2015-01-02      TSLA  26.123333  25.620192  26.123333  11.873179  49.228223
2015-01-02     GOOGL  26.123333  25.620192  26.123333  11.873179  49.228223
2015-01-02     MSFT  26.123333  25.620192  26.123333  11.873179  49.228223
2015-01-02     AMZN  26.123333  25.620192  26.123333  11.873179  49.228223

```

```

      BB_Lower      RSI_14  Log_Return      OBV  Prev_Close \
Date
2015-01-02  1.735509  46.167054  -0.625640      0.0  27.332500
2015-01-02  1.735509  46.167054  -0.625640 -71466000.0  27.332500
2015-01-02  1.735509  46.167054   0.593859 -44986000.0  14.620667
2015-01-02  1.735509  46.167054   0.568733 -17072100.0  26.477501
2015-01-02  1.735509  46.167054  -1.108974 -72736100.0  46.759998

```

```

      Prev_Volume  DayOfWeek  Month
Date
2015-01-02  212818400.0      4      1
2015-01-02  212818400.0      4      1
2015-01-02   71466000.0      4      1
2015-01-02  26480000.0      4      1
2015-01-02  27913900.0      4      1

```

### 1.4.3 3. Filter Company Data

```

[40]: tweets_filtered = pd.read_csv("../data/filtered/tweets_filtered.csv")

```

```

[41]: company_list = ['AAPL', 'AMZN', 'MSFT', 'GOOGL', 'TSLA']

```

```
[42]: # Filter for selected stocks
filter_tweets = tweets_filtered[tweets_filtered['Stock Name'].
    ↪isin(company_list)].copy()

# Group and compute daily average sentiment
daily_sentiment = (
    filter_tweets
    .groupby(['Stock Name', 'Date'])['Vader_Polarity']
    .mean()
    .reset_index()
)

# Set Date as index (for plotting or merging)
daily_sentiment['Date'] = pd.to_datetime(daily_sentiment['Date'])
daily_sentiment = daily_sentiment.set_index('Date').sort_index()

daily_sentiment.head()
```

```
[42]:      Stock Name  Vader_Polarity
Date
2015-01-02      AAPL      0.276950
2015-01-02      GOOGL      0.470091
2015-01-02      TSLA      0.227755
2015-01-02      AMZN      0.222096
2015-01-02      MSFT      0.260802
```

```
[43]: filter_stocks = stocks[stocks['Stock Name'].isin(company_list)]
filter_stocks.head()
```

```
[43]:      Open      High      Low      Close  Adj Close      Volume \
Date
2015-01-02  27.847500  27.860001  26.837500  27.332500  24.320429  212818400
2015-01-02  14.858000  14.883333  14.217333  14.620667  14.620667   71466000
2015-01-02  26.629999  26.790001  26.393999  26.477501  26.351515   26480000
2015-01-02  46.660000  47.419998  46.540001  46.759998  40.072136   27913900
2015-01-02  15.629000  15.737500  15.348000  15.426000  15.426000   55664000

      Stock Name      SMA_5      SMA_20      BB_Mid      BB_Std      BB_Upper \
Date
2015-01-02      AAPL  26.123333  25.620192  26.123333  11.873179  49.228223
2015-01-02      TSLA  26.123333  25.620192  26.123333  11.873179  49.228223
2015-01-02      GOOGL  26.123333  25.620192  26.123333  11.873179  49.228223
2015-01-02      MSFT  26.123333  25.620192  26.123333  11.873179  49.228223
2015-01-02      AMZN  26.123333  25.620192  26.123333  11.873179  49.228223

      BB_Lower      RSI_14  Log_Return      OBV  Prev_Close \
Date
```

2015-01-02	1.735509	46.167054	-0.625640	0.0	27.332500
2015-01-02	1.735509	46.167054	-0.625640	-71466000.0	27.332500
2015-01-02	1.735509	46.167054	0.593859	-44986000.0	14.620667
2015-01-02	1.735509	46.167054	0.568733	-17072100.0	26.477501
2015-01-02	1.735509	46.167054	-1.108974	-72736100.0	46.759998

Date	Prev_Volume	DayOfWeek	Month
2015-01-02	212818400.0	4	1
2015-01-02	212818400.0	4	1
2015-01-02	71466000.0	4	1
2015-01-02	26480000.0	4	1
2015-01-02	27913900.0	4	1

#### 1.4.4 4. Merging (Stock + Sentiment)

```
[44]: def merging(company_list, filter_stocks, daily_sentiment):
    """
    Create a joined dataset of stock data and sentiment scores for a list of
    ↪stock tickers.

    Parameters:
    -----
    company_list : list
        List of stock ticker symbols (e.g., ['AAPL', 'MSFT', 'GOOGL'])
    filter_stocks : pandas.DataFrame
        DataFrame containing stock price data with 'Stock Name' column
    daily_sentiment : pandas.DataFrame
        DataFrame containing sentiment scores with 'Stock Name' column

    Returns:
    -----
    dict
        Dictionary with ticker symbols as keys and joined DataFrames as values
    """
    # Dictionary to store results
    result_dict = {}

    for company in company_list:
        stock = filter_stocks[filter_stocks['Stock Name'] == company].copy()
        ticker_sentiment = daily_sentiment[daily_sentiment['Stock Name'] ==
        ↪company].copy()

        # Normalize datetime index (remove time, timezone)
        stock.index = pd.to_datetime(stock.index).normalize()
        ticker_sentiment.index = pd.to_datetime(ticker_sentiment.index).
        ↪normalize()
```

```

# Perform inner join to keep only dates present in both
joined_data = stock.join(ticker_sentiment[['Vader_Polarity']],
↳how='inner')

result_dict[company] = joined_data

return result_dict

stock_data_dict = merging(company_list, filter_stocks, daily_sentiment)

```

```

[45]: AAPL = stock_data_dict['AAPL']

AAPL.head()

```

```

[45]:
      Open      High      Low      Close  Adj Close      Volume \
Date
2015-01-02  27.847500  27.860001  26.837500  27.332500  24.320429  212818400
2015-01-05  27.072500  27.162500  26.352501  26.562500  23.635286  257142000
2015-01-06  26.635000  26.857500  26.157499  26.565001  23.637505  263188400
2015-01-07  26.799999  27.049999  26.674999  26.937500  23.968958  160423600
2015-01-08  27.307501  28.037500  27.174999  27.972500  24.889906  237458000

```

```

      Stock Name      SMA_5      SMA_20      BB_Mid  ...  BB_Upper \
Date
2015-01-02      AAPL  26.123333  25.620192  26.123333  ...  49.228223
2015-01-05      AAPL  19.415400  25.620192  19.415400  ...  49.228223
2015-01-06      AAPL  31.822701  25.620192  31.822701  ...  49.228223
2015-01-07      AAPL  27.564067  25.620192  27.564067  ...  49.228223
2015-01-08      AAPL  19.783466  25.672308  19.783466  ...  43.502581

```

```

      BB_Lower      RSI_14  Log_Return      OBV  Prev_Close \
Date
2015-01-02  1.735509  46.167054  -0.625640      0.0  27.332500
2015-01-05  1.735509  46.167054   0.640015  89576400.0  14.006000
2015-01-06  1.735509  46.167054  -0.541409 -170386000.0  45.650002
2015-01-07  1.735509  53.500699   0.649948 -41384800.0  14.063333
2015-01-08 -3.935649  54.182309   0.621640 205035200.0  15.023000

```

```

      Prev_Volume  DayOfWeek  Month  Vader_Polarity
Date
2015-01-02  212818400.0      4      1      0.276950
2015-01-05  80527500.0      0      1      0.251420
2015-01-06  36447900.0      1      1      0.281760
2015-01-07  44526000.0      2      1      0.271373
2015-01-08  61768000.0      3      1      0.321796

```

[5 rows x 21 columns]

```
[46]: for company in company_list:
      df = stock_data_dict[company]
      df.to_csv(f"./data/filtered/{company}_filtered.csv", index=True)
```

## 1.4.5 5. Visualizing Related Features

```
[47]: company_list = ['TSLA', 'AAPL', 'AMZN', 'GOOGL']
      stock_data_dict = {}

      for symbol in company_list:
          path = f"./data/filtered/{symbol}_filtered.csv"
          stock_data_dict[symbol] = pd.read_csv(path)
```

```
[48]: tweets_filtered = pd.read_csv("./data/filtered/tweets_filtered.csv")
```

```
[49]: TSLA = stock_data_dict['TSLA']
      TSLA.head()
```

```
[49]:
```

	Date	Open	High	Low	Close	Adj Close	\
0	2015-01-02	14.858000	14.883333	14.217333	14.620667	14.620667	
1	2015-01-05	14.303333	14.433333	13.810667	14.006000	14.006000	
2	2015-01-06	14.004000	14.280000	13.614000	14.085333	14.085333	
3	2015-01-07	14.223333	14.318667	13.985333	14.063333	14.063333	
4	2015-01-08	14.187333	14.253333	14.000667	14.041333	14.041333	

	Volume	Stock	Name	SMA_5	SMA_20	...	BB_Upper	BB_Lower	\
0	71466000	TSLA		26.123333	25.620192	...	49.228223	1.735509	
1	80527500	TSLA		23.454899	25.620192	...	49.228223	1.735509	
2	93928500	TSLA		31.838567	25.620192	...	49.228223	1.735509	
3	44526000	TSLA		25.129466	25.620192	...	49.228223	1.735509	
4	51637500	TSLA		25.994467	25.587975	...	50.001643	1.987290	

	RSI_14	Log_Return	OBV	Prev_Close	Prev_Volume	DayOfWeek	\
0	46.167054	-0.625640	-71466000.0	27.332500	212818400.0	4	
1	46.167054	-0.075838	-167565600.0	15.109500	55484000.0	0	
2	46.167054	-0.634461	-264314500.0	26.565001	263188400.0	1	
3	41.060477	-1.190058	-201808400.0	46.230000	29114100.0	2	
4	41.785148	-1.220617	109988900.0	47.590000	29645200.0	3	

	Month	Vader_Polarity
0	1	0.227755
1	1	0.366453
2	1	0.202426
3	1	0.251155
4	1	0.279833

[5 rows x 22 columns]

```
[50]: # Combine all stock DataFrames into one
stocks = pd.concat(stock_data_dict.values(), ignore_index=True)

# datetime
stocks['Date'] = pd.to_datetime(stocks['Date'])
stocks = stocks.sort_values(['Date', 'Stock Name']).reset_index(drop=True)

print(stocks['Stock Name'].value_counts())
```

```
Stock Name
AMZN      1258
GOOGL     1258
TSLA      1258
AAPL      1255
Name: count, dtype: int64
```

```
[51]: stocks.head()
```

```
[51]:
```

	Date	Open	High	Low	Close	Adj Close	\
0	2015-01-02	27.847500	27.860001	26.837500	27.332500	24.320429	
1	2015-01-02	15.629000	15.737500	15.348000	15.426000	15.426000	
2	2015-01-02	26.629999	26.790001	26.393999	26.477501	26.351515	
3	2015-01-02	14.858000	14.883333	14.217333	14.620667	14.620667	
4	2015-01-05	27.072500	27.162500	26.352501	26.562500	23.635286	

	Volume	Stock Name	SMA_5	SMA_20	...	BB_Upper	BB_Lower	\
0	212818400	AAPL	26.123333	25.620192	...	49.228223	1.735509	
1	55664000	AMZN	26.123333	25.620192	...	49.228223	1.735509	
2	26480000	GOOGL	26.123333	25.620192	...	49.228223	1.735509	
3	71466000	TSLA	26.123333	25.620192	...	49.228223	1.735509	
4	257142000	AAPL	19.415400	25.620192	...	49.228223	1.735509	

	RSI_14	Log_Return	OBV	Prev_Close	Prev_Volume	DayOfWeek	\
0	46.167054	-0.625640	0.0	27.332500	212818400.0	4	
1	46.167054	-1.108974	-72736100.0	46.759998	27913900.0	4	
2	46.167054	0.593859	-44986000.0	14.620667	71466000.0	4	
3	46.167054	-0.625640	-71466000.0	27.332500	212818400.0	4	
4	46.167054	0.640015	89576400.0	14.006000	80527500.0	0	

	Month	Vader_Polarity
0	1	0.276950
1	1	0.222096
2	1	0.470091
3	1	0.227755
4	1	0.251420

[5 rows x 22 columns]

```
[52]: stocks['Date'].unique().value_counts()
```

```
[52]: 2015-01-02    1
      2015-01-05    1
      2015-01-06    1
      2015-01-07    1
      2015-01-08    1
      ..
      2019-12-24    1
      2019-12-26    1
      2019-12-27    1
      2019-12-30    1
      2019-12-31    1
      Name: count, Length: 1258, dtype: int64
```

```
[53]: def plot_price_sma_grid(df_all, company_list):
      fig, axs = plt.subplots(2, 2, figsize=(21, 12)) # 14x6 per plot
      axs = axs.flatten()

      for i, company in enumerate(company_list):
          df = df_all[df_all['Stock Name'] == company]
          axs[i].plot(df['Date'], df['Close'], label='Close', color='black')
          axs[i].plot(df['Date'], df['SMA_5'], label='SMA 5', color='red')
          axs[i].set_title(f'{company} - Closing Price vs SMA 5')
          axs[i].legend()
          axs[i].grid(True)

      plt.tight_layout()
      plt.show()

      def plot_log_return_grid(df_all, company_list):
          fig, axs = plt.subplots(2, 2, figsize=(21, 12))
          axs = axs.flatten()

          for i, company in enumerate(company_list):
              df = df_all[df_all['Stock Name'] == company]
              axs[i].plot(df['Date'], df['Log_Return'], color='orange')
              axs[i].set_title(f'{company} - Daily Log Return')
              axs[i].grid(True)

          plt.tight_layout()
          plt.show()

      def plot_rsi_grid(df_all, company_list):
```

```

fig, axs = plt.subplots(2, 2, figsize=(21, 12))
axs = axs.flatten()

for i, company in enumerate(company_list):
    df = df_all[df_all['Stock Name'] == company]
    axs[i].plot(df['Date'], df['RSI_14'], color='purple')
    axs[i].axhline(70, color='red', linestyle='--')
    axs[i].axhline(30, color='green', linestyle='--')
    axs[i].set_title(f'{company} - RSI (14-day)')
    axs[i].grid(True)

plt.tight_layout()
plt.show()


def plot_volume_grid(df_all, company_list):
    fig, axs = plt.subplots(2, 2, figsize=(21, 12))
    axs = axs.flatten()

    for i, company in enumerate(company_list):
        df = df_all[df_all['Stock Name'] == company]
        axs[i].plot(df['Date'], df['Volume'], color='gray')
        axs[i].set_title(f'{company} - Daily Trading Volume')
        axs[i].grid(True)

    plt.tight_layout()
    plt.show()


```

[54]: `plot_price_sma_grid(stocks, company_list)`

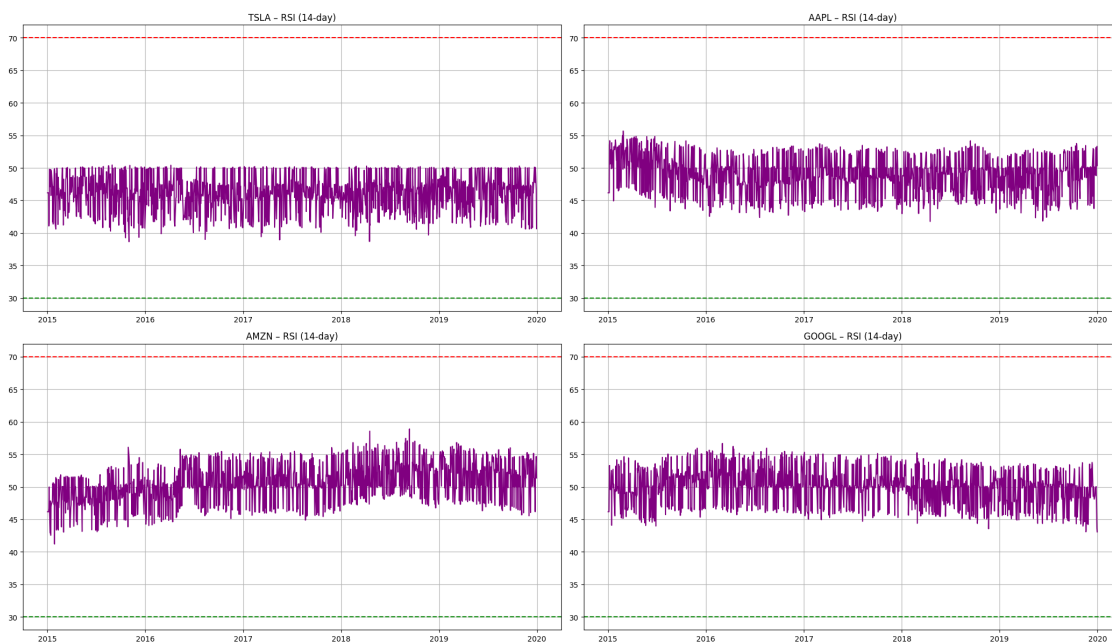




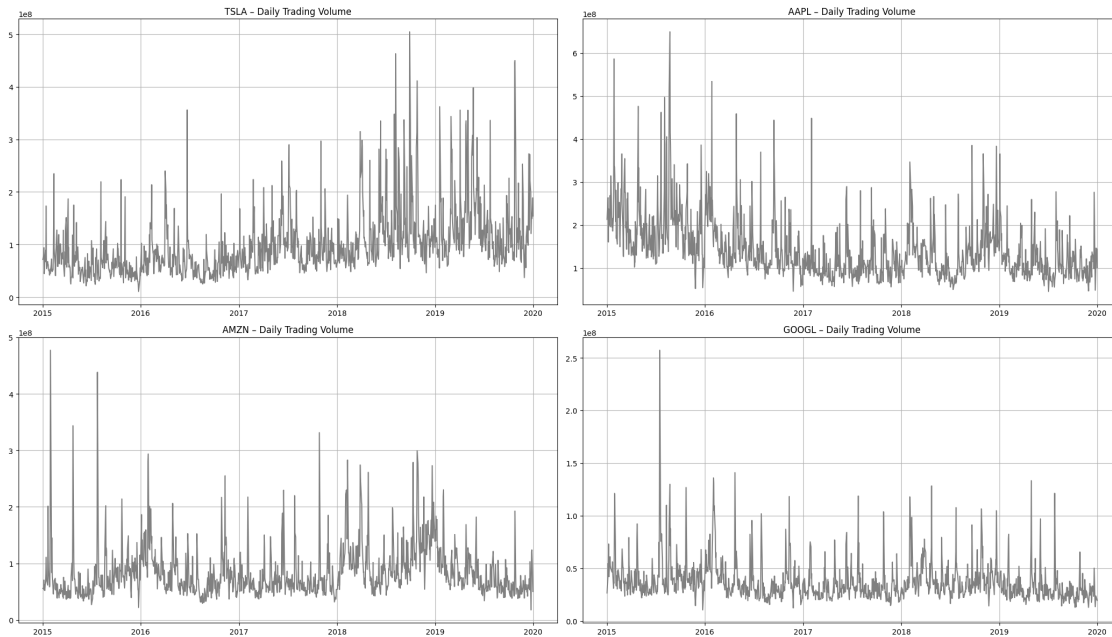
```
[55]: plot_log_return_grid(stocks, company_list)
```



```
[56]: plot_rsi_grid(stocks, company_list)
```



```
[57]: plot_volume_grid(stocks, company_list)
```



```
[58]: def plot_price_vs_sentiment(df_all, company_list,
    ↪ sentiment_col='Vader_Polarity'):
    fig, axs = plt.subplots(2, 2, figsize=(21, 12)) # Grid for 4 companies
    axs = axs.flatten()

    for i, company in enumerate(company_list):
        df = df_all[df_all['Stock Name'] == company]
        ax1 = axs[i]

        # Left y-axis: Close Price
        ax1.plot(df['Date'], df['Close'], color='black', linestyle='--',
    ↪ label='Close')
        ax1.set_ylabel('Close Price ($)', color='black', fontsize=12)
        ax1.tick_params(axis='y', labelcolor='black')
        ax1.grid(True) #

        # Twin y-axis for sentiment
        ax2 = ax1.twinx()
        ax2.plot(df['Date'], df[sentiment_col], color='orange',
    ↪ label=sentiment_col)
        ax2.set_ylabel(sentiment_col, color='orange', fontsize=12)
        ax2.tick_params(axis='y', labelcolor='orange')
        ax2.grid(False) #
```

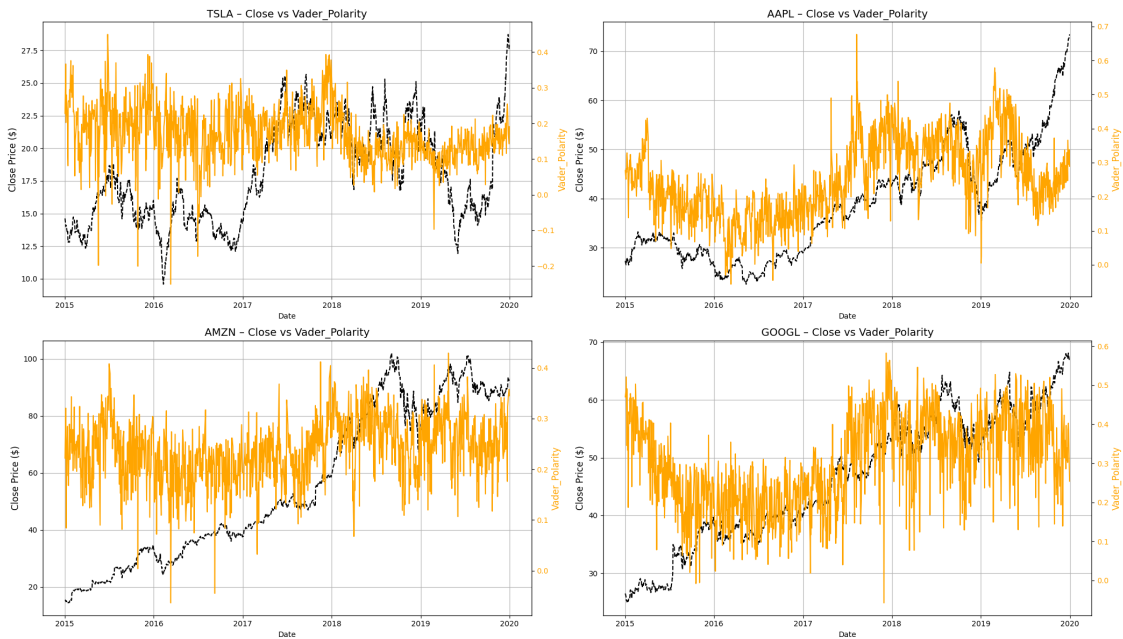
```

# Title and grid
ax1.set_title(f'{company} - Close vs {sentiment_col}', fontsize=14)
ax1.set_xlabel('Date')

plt.tight_layout()
plt.show()

```

```
[59]: plot_price_vs_sentiment(stocks, company_list, sentiment_col='Vader_Polarity')
```



## 1.5 Feature Selection

```

[60]: TSLA = stock_data_dict['TSLA']
df = TSLA.copy()

df.head()

```

```

[60]:
      Date      Open      High      Low      Close  Adj Close  \
0  2015-01-02  14.858000  14.883333  14.217333  14.620667  14.620667
1  2015-01-05  14.303333  14.433333  13.810667  14.006000  14.006000
2  2015-01-06  14.004000  14.280000  13.614000  14.085333  14.085333
3  2015-01-07  14.223333  14.318667  13.985333  14.063333  14.063333
4  2015-01-08  14.187333  14.253333  14.000667  14.041333  14.041333

      Volume Stock Name      SMA_5      SMA_20  ...  BB_Upper  BB_Lower  \
0  71466000      TSLA  26.123333  25.620192  ...  49.228223  1.735509
1  80527500      TSLA  23.454899  25.620192  ...  49.228223  1.735509

```

2	93928500	TSLA	31.838567	25.620192	...	49.228223	1.735509
3	44526000	TSLA	25.129466	25.620192	...	49.228223	1.735509
4	51637500	TSLA	25.994467	25.587975	...	50.001643	1.987290

	RSI_14	Log_Return	OBV	Prev_Close	Prev_Volume	DayOfWeek	\
0	46.167054	-0.625640	-71466000.0	27.332500	212818400.0		4
1	46.167054	-0.075838	-167565600.0	15.109500	55484000.0		0
2	46.167054	-0.634461	-264314500.0	26.565001	263188400.0		1
3	41.060477	-1.190058	-201808400.0	46.230000	29114100.0		2
4	41.785148	-1.220617	109988900.0	47.590000	29645200.0		3

	Month	Vader_Polarity
0	1	0.227755
1	1	0.366453
2	1	0.202426
3	1	0.251155
4	1	0.279833

[5 rows x 22 columns]

```
[61]: df.columns
```

```
[61]: Index(['Date', 'Open', 'High', 'Low', 'Close', 'Adj Close', 'Volume',
          'Stock Name', 'SMA_5', 'SMA_20', 'BB_Mid', 'BB_Std', 'BB_Upper',
          'BB_Lower', 'RSI_14', 'Log_Return', 'OBV', 'Prev_Close', 'Prev_Volume',
          'DayOfWeek', 'Month', 'Vader_Polarity'],
          dtype='object')
```

### 1.5.1 1. Avoid Data Leakage

```
[62]: # Lag technical indicators to avoid leakage
lag_cols = [
    'Adj Close', 'High', 'Low', 'Volume', 'SMA_5', 'SMA_20',
    'BB_Mid', 'BB_Std', 'BB_Upper', 'BB_Lower',
    'RSI_14', 'Log_Return', 'OBV', 'Vader_Polarity'
]

for col in lag_cols:
    if col == "Adj Close":
        df[f"{col} (lag1)"] = df[col].shift(1)
    else:
        df[col] = df[col].shift(1)

feature_cols = [
    'Adj Close',
    'Open', 'High', 'Low', 'Volume', 'Adj Close (lag1)',
    'SMA_5', 'SMA_20', 'BB_Mid', 'BB_Upper', 'BB_Lower',
```

```

    'OBV', 'Log_Return', 'DayOfWeek', 'Month',
    'Vader_Polarity'
]

df = df[feature_cols]

```

```

[63]: feature_cols = [
    'Adj Close',          # Adj Close
    'Open', 'High', 'Low', 'Volume',
    'Adj Close (lag1)',    #
    'SMA_5', 'SMA_20', 'BB_Mid', 'BB_Upper', 'BB_Lower',
    'OBV', 'Log_Return', 'DayOfWeek', 'Month',
    'Vader_Polarity'
]

target_col = 'Adj Close' # or 'Log_Return'

```

## 1.5.2 2. Feature Importance

```


[64]: # --- X_corr Adj Close shift ---
X_corr = df[feature_cols].copy()
X_corr = X_corr.iloc[:-1, :] # label y = shift(-1)

# --- 2. Correlation Heatmap ---
plt.figure(figsize=(14, 11))
corr_matrix = X_corr.corr()
sns.heatmap(corr_matrix, annot=True, fmt=".2f", cmap='coolwarm',
            annot_kws={"size": 10}) #

plt.title("Feature Correlation Matrix", fontsize=16) #
plt.xticks(rotation=45, ha='right', fontsize=12) # x
plt.yticks(rotation=0, fontsize=12) # y

plt.tight_layout()
plt.show()

```



```
[65]: # 1. X Adj Close y
X = df[feature_cols].copy()
X = X.drop(columns=['Adj Close']) # Adj Close
y = df['Adj Close'] # y

# 2. lag SMA
X = X.iloc[1:, :] # lag1
y = y.iloc[1:] #

# 3.
scaler = StandardScaler()
X_scaled = pd.DataFrame(scaler.fit_transform(X), columns=X.columns)

# 4.
X_train, X_val, y_train, y_val = train_test_split(
    X_scaled, y, test_size=0.3, shuffle=False
)
```

```
[66]: # --- 3. XGBoost Feature Importance ---
xgb_model = xgb.XGBRegressor(n_estimators=100)
xgb_model.fit(X_train, y_train)


importances = xgb_model.get_booster().get_score(importance_type='gain')
importances = sorted(importances.items(), key=lambda x: x[1], reverse=True)

imp_df = pd.DataFrame(importances, columns=['Feature', 'Importance']).head(10)

plt.figure(figsize=(10, 5))
sns.barplot(
    data=imp_df,
    x='Importance',
    y='Feature',
    hue='Feature',
    palette='viridis',
    dodge=False,
    legend=False
)


# x
for i, v in enumerate(imp_df['Importance']):
    plt.text(v + 0.5, i, f'{v:.2f}', va='center', fontsize=10)

plt.title("Top 10 XGBoost Feature Importance (Gain)", fontsize=14)
plt.xlabel("Importance Score", fontsize=12)
plt.ylabel("")
plt.xticks(fontsize=10)
plt.yticks(fontsize=10)
plt.tight_layout()
plt.show()
```



```
[67]: # --- 4. SHAP with PermutationExplainer (CPU-friendly) ---
explainer = shap.Explainer(xgb_model.predict, X_val, algorithm="permutation")
shap_values = explainer(X_val)

#      + 10
plt.gcf().set_size_inches(10, 5)
shap.plots.bar(shap_values, max_display=10)
plt.tight_layout()
plt.show()
```



<Figure size 640x480 with 0 Axes>

```
[68]: # ===== 1. =====
X_corr_check = df[feature_cols].copy().iloc[: -1] #
corr_matrix = X_corr_check.corr().abs()
upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).astype(bool))
to_drop_corr = [col for col in upper.columns if any(upper[col] > 0.90)]

print(" Dropped (correlated > 0.90):")
for col in to_drop_corr:
    print(f"    · {col}")
```



```
# ===== 2. XGBoost Top 10 =====
# xgb_model
importances_dict = xgb_model.get_booster().get_score(importance_type='gain')
top10_sorted = sorted(importances_dict.items(), key=lambda x: x[1],
    ↪reverse=True)[:10]

print("\n Top 10 XGBoost Features (by Gain):")
for i, (feat, score) in enumerate(top10_sorted, 1):
    print(f" {i:>2d}. {feat:<20} → {score:.2f}")
```

Dropped (correlated > 0.90):

- Open
- High
- Low
- Adj Close (lag1)
- SMA\_20
- BB\_Mid
- BB\_Upper
- OBV

Top 10 XGBoost Features (by Gain):

- |                     |         |
|---------------------|---------|
| 1. Open             | → 34.80 |
| 2. Adj Close (lag1) | → 5.03  |
| 3. High             | → 0.61  |
| 4. Low              | → 0.59  |
| 5. OBV              | → 0.26  |
| 6. SMA_20           | → 0.09  |
| 7. Month            | → 0.08  |
| 8. BB_Upper         | → 0.07  |
| 9. Log_Return       | → 0.06  |
| 10. SMA_5           | → 0.06  |

## 1.6 Time Series Train-Test Split

### 1.6.1 1. Loading Pre-processed Data

```
[133]: company_list = ['TSLA', 'AAPL', 'AMZN', 'GOOGL', 'MSFT']
stock_data_dict = {}

for symbol in company_list:
    path = f"./data/filtered/{symbol}_filtered.csv"
    stock_data_dict[symbol] = pd.read_csv(path)
```

```
[134]: TSLA = stock_data_dict['TSLA']
AAPL = stock_data_dict['AAPL']
AMZN = stock_data_dict['AMZN']
GOOGL = stock_data_dict['GOOGL']
```

```
MSFT = stock_data_dict['MSFT']
```

```
[135]: df = TSLA.copy()
df.head()
```

```
[135]:
```

	Date	Open	High	Low	Close	Adj Close	\
0	2015-01-02	14.858000	14.883333	14.217333	14.620667	14.620667	
1	2015-01-05	14.303333	14.433333	13.810667	14.006000	14.006000	
2	2015-01-06	14.004000	14.280000	13.614000	14.085333	14.085333	
3	2015-01-07	14.223333	14.318667	13.985333	14.063333	14.063333	
4	2015-01-08	14.187333	14.253333	14.000667	14.041333	14.041333	

	Volume	Stock Name	SMA_5	SMA_20	...	BB_Upper	BB_Lower	\
0	71466000	TSLA	26.123333	25.620192	...	49.228223	1.735509	
1	80527500	TSLA	23.454899	25.620192	...	49.228223	1.735509	
2	93928500	TSLA	31.838567	25.620192	...	49.228223	1.735509	
3	44526000	TSLA	25.129466	25.620192	...	49.228223	1.735509	
4	51637500	TSLA	25.994467	25.587975	...	50.001643	1.987290	

	RSI_14	Log_Return	OBV	Prev_Close	Prev_Volume	DayOfWeek	\
0	46.167054	-0.625640	-71466000.0	27.332500	212818400.0	4	
1	46.167054	-0.075838	-167565600.0	15.109500	55484000.0	0	
2	46.167054	-0.634461	-264314500.0	26.565001	263188400.0	1	
3	41.060477	-1.190058	-201808400.0	46.230000	29114100.0	2	
4	41.785148	-1.220617	109988900.0	47.590000	29645200.0	3	

	Month	Vader_Polarity
0	1	0.227755
1	1	0.366453
2	1	0.202426
3	1	0.251155
4	1	0.279833

```
[5 rows x 22 columns]
```

## 1.6.2 2. Avoid Data Leakage

```
[136]: # Lag technical indicators to avoid leakage
lag_cols = [
    'Adj Close', 'High', 'Low', 'Volume', 'SMA_5', 'SMA_20',
    'BB_Mid', 'BB_Std', 'BB_Upper', 'BB_Lower',
    'RSI_14', 'Log_Return', 'OBV', 'Vader_Polarity'
]

for col in lag_cols:
    if col == "Adj Close":
        df[f"{col} (lag1)"] = df[col].shift(1)
```

```

else:
    df[col] = df[col].shift(1)

feature_cols = [
    'Adj Close',
    'Open', 'High', 'Low', 'Volume', 'Adj Close (lag1)',
    'SMA_5', 'SMA_20', 'BB_Mid', 'BB_Upper', 'BB_Lower',
    'OBV', 'Log_Return', 'DayOfWeek', 'Month',
    'Vader_Polarity'
]

df = df[feature_cols]

```

```
[137]: df.head()
```

```

[137]:
   Adj Close      Open      High      Low      Volume  Adj Close (lag1) \
0  14.620667  14.858000      NaN      NaN      NaN      NaN
1  14.006000  14.303333  14.883333  14.217333  71466000.0    14.620667
2  14.085333  14.004000  14.433333  13.810667  80527500.0    14.006000
3  14.063333  14.223333  14.280000  13.614000  93928500.0    14.085333
4  14.041333  14.187333  14.318667  13.985333  44526000.0    14.063333

      SMA_5      SMA_20      BB_Mid      BB_Upper      BB_Lower      OBV \
0      NaN      NaN      NaN      NaN      NaN      NaN
1  26.123333  25.620192  26.123333  49.228223  1.735509 -71466000.0
2  23.454899  25.620192  23.454899  49.228223  1.735509 -167565600.0
3  31.838567  25.620192  31.838567  49.228223  1.735509 -264314500.0
4  25.129466  25.620192  25.129466  49.228223  1.735509 -201808400.0

      Log_Return      DayOfWeek      Month      Vader_Polarity
0      NaN          4          1          NaN
1  -0.625640          0          1      0.227755
2  -0.075838          1          1      0.366453
3  -0.634461          2          1      0.202426
4  -1.190058          3          1      0.251155

```

```

[138]: # feature_cols = [
#     'Adj Close', # should be the first one for Y
#     'Adj Close (lag1)',
#     'SMA_5',          # short-term trend
#     'Volume',
#     'BB_Mid',          # risk signal
#     'Log_Return',
#     'DayOfWeek',
#     'Month',
#     'Vader_Polarity' # should be the last one for SENTIMENT
# ]

```

```
[139]: feature_cols = [
        'Adj Close', # should be the first one for Y
        'Adj Close (lag1)',
        'SMA_5',      # short-term trend
        'Volume',
        'BB_Mid',     # risk signal
        'Log_Return',
        'DayOfWeek',
        'Month',
        'Vader_Polarity' # should be the last one for SENTIMENT
    ]
```

### 1.6.3 Define Rolling Window & Prediction Day

```
[140]: # Step 0: Define sliding window parameters
n_past = 5 # Use past 5 days
n_future = 1 # Predict next 1 day
```

### 1.6.4 Train-Test Split

```
[141]: train_size = 0.7
train_split_idx = int(train_size * len(df))
```

```
[142]: df_filtered = df[feature_cols]
df_filtered = df_filtered.iloc[1:] # delete nan lag1

# Step 0: Define split boundaries BEFORE scaling
train_df = df_filtered.iloc[:train_split_idx]
test_df = df_filtered.iloc[train_split_idx:]

# Step 1: Fit scaler only on training data (Avoid Data Leakage)
scaler = MinMaxScaler()
scaler.fit(train_df)

# Step 2: Scale training and test data separately
train_scaled = scaler.transform(train_df)
test_scaled = scaler.transform(test_df)

# Step 3: For inference later, only scale ['Adj Close']
scaler_for_inference = MinMaxScaler()
actual_scaled_close = scaler_for_inference.fit_transform(
    df_filtered[['Adj Close']]
)

# Step 3: Reconstruct sliding windows for train and test
def create_sequences(data, n_past, n_future):
    X, y = [], []
```

```

for i in range(n_past, len(data) - n_future + 1):
    X.append(data[i - n_past:i, 1:])
    y.append(data[i + n_future - 1:i + n_future, [0]]) # Predict Adj Close
return np.array(X), np.array(y)

trainX, trainY = create_sequences(train_scaled, n_past, n_future)
testX, testY = create_sequences(test_scaled, n_past, n_future)

# trainY = trainY.reshape(-1, 1)
# testY = testY.reshape(-1, 1)

print('TrainX shape = {}'.format(trainX.shape))
print('TrainY shape = {}'.format(trainY.shape))
print('TestX shape = {}'.format(testX.shape))
print('TestY shape = {}'.format(testY.shape))

```

```

TrainX shape = (875, 5, 8)
TrainY shape = (875, 1, 1)
TestX shape = (372, 5, 8)
TestY shape = (372, 1, 1)

```

### 1.6.5 (Un)Sentiment Split

```

[143]: # Without Sentiment (Baseline Model)
trainX_wo_tweet = trainX[:, :, :-1] # Exclude last feature
testX_wo_tweet = testX[:, :, :-1]
trainY_wo_tweet = trainY
testY_wo_tweet = testY

# With Sentiment (Tweet-based Model)
trainX_with_tweet = trainX
testX_with_tweet = testX
trainY_with_tweet = trainY
testY_with_tweet = testY

```

```

[144]: print(trainX_with_tweet.shape, trainY_with_tweet.shape)

```

```

(875, 5, 8) (875, 1, 1)

```

## 2 Modeling

### 2.0.1 Random Seed

```

[145]: def set_all_seeds(seed=42):
        os.environ['PYTHONHASHSEED'] = str(seed)
        random.seed(seed)
        np.random.seed(seed)

```

```

tf.random.set_seed(seed)
tf.config.experimental.enable_op_determinism() # TensorFlow 2.12+

torch.manual_seed(seed)
if torch.cuda.is_available():
    torch.cuda.manual_seed_all(seed)
    torch.backends.cudnn.deterministic = True
    torch.backends.cudnn.benchmark = False

set_all_seeds(seed=42)

```

## 2.0.2 CNN + BiLSTM

### Pre-load Plotting Functions

```

[146]: def plot_metrics(model, history, X_train, y_train, X_val, y_val,
    ↪ sentiment_mode):
    """
        RMSE    MAE        RMSE / MAE
    """
    history_data = history.history #

    # 1.
    y_train_pred = model.predict(X_train, verbose=0).squeeze()
    y_val_pred    = model.predict(X_val, verbose=0).squeeze()

    y_train_true = y_train.squeeze()
    y_val_true   = y_val.squeeze()
    # y_train_true = y_train.reshape(-1, 1)
    # y_val_true   = y_val.reshape(-1, 1)

    # 2.
    train_rmse = np.sqrt(mean_squared_error(y_train_true, y_train_pred))
    val_rmse    = np.sqrt(mean_squared_error(y_val_true, y_val_pred))

    train_mae = mean_absolute_error(y_train_true, y_train_pred)
    val_mae    = mean_absolute_error(y_val_true, y_val_pred)

    title_suffix = " w/ Tweets" if sentiment_mode == "with" else "w/o Tweets"

    # 3.
    print(f" {title_suffix} RMSE: Train = {train_rmse:.4f}, Test = {val_rmse:.4f}")
    ↪4f}")
    print(f" {title_suffix} MAE : Train = {train_mae:.4f}, Test = {val_mae:.4f}")
    ↪4f}")

    # 4.    loss
    plt.figure(figsize=(21, 9))

```

```

# ---- RMSE subplot
plt.subplot(1, 2, 1)
plt.plot(np.sqrt(history_data['loss']), label='Train RMSE')
plt.plot(np.sqrt(history_data['val_loss']), label='Val RMSE')
plt.title(f'RMSE - {title_suffix}', fontsize=18)
plt.xlabel('Epoch', fontsize=18);
plt.ylabel('RMSE', fontsize=18)
plt.xticks(fontsize=16)
plt.yticks(fontsize=16)
plt.legend()

# ---- MAE subplot
plt.subplot(1, 2, 2)
if 'mae' in history_data:
    plt.plot(history_data['mae'], label='Train MAE')
    plt.plot(history_data['val_mae'], label='Val MAE')
else:
    # fallback: flat line (not recommended long term)
    plt.plot([train_mae] * len(history_data['loss']), label='Train MAE_␣
↪(flat)')
    plt.plot([val_mae] * len(history_data['val_loss']), label='Val MAE_␣
↪(flat)')
    plt.title(f'MAE - {title_suffix}', fontsize=18)
    plt.xlabel('Epoch', fontsize=18);
    plt.ylabel('MAE', fontsize=18)
    plt.xticks(fontsize=16)
    plt.yticks(fontsize=16)
    plt.legend()

# Add supertitle here
plt.suptitle(f'Model Training Metrics - CNN+BiLSTM', fontsize=24)
plt.tight_layout()

plt.show()

```

```

[147]: # Helper Function: get predicted and true values in original (unscaled) price_␣
↪units
def inv_preds(model, X, y):
    # Predict using the model and reshape to [samples, output_dim]
    p = model.predict(X, verbose=0).reshape(-1, y.shape[-1])

    # Reshape ground truth to same shape for inverse transformation
    g = y.reshape(-1, y.shape[-1])

    # Apply inverse scaling to recover original price scale
    # inv_p = scaler_for_inference.inverse_transform(p)

```

```

# inv_g = scaler_for_inference.inverse_transform(g)

# Return only the first column ('Adj Close') from each
# return inv_p[:, 0], inv_g[:, 0]

# Use the full-feature scaler
inv_p = scaler.inverse_transform(
    np.concatenate([p, np.zeros((p.shape[0], scaler.n_features_in_ - 1))],
axis=1)
[:, 0]
inv_g = scaler.inverse_transform(
    np.concatenate([g, np.zeros((g.shape[0], scaler.n_features_in_ - 1))],
axis=1)
[:, 0]

return inv_p, inv_g

```

```
[148]: df_filtered.head()
```

```
[148]:
```

	Adj Close	Adj Close (lag1)	SMA_5	Volume	BB_Mid	Log_Return	\
1	14.006000	14.620667	26.123333	71466000.0	26.123333	-0.625640	
2	14.085333	14.006000	23.454899	80527500.0	23.454899	-0.075838	
3	14.063333	14.085333	31.838567	93928500.0	31.838567	-0.634461	
4	14.041333	14.063333	25.129466	44526000.0	25.129466	-1.190058	
5	13.777333	14.041333	25.994467	51637500.0	25.994467	-1.220617	

	DayOfWeek	Month	Vader_Polarity
1	0	1	0.227755
2	1	1	0.366453
3	2	1	0.202426
4	3	1	0.251155
5	4	1	0.279833

```
[149]: def plot_pred_vs_actual(
    model, X_train, y_train, X_test, y_test, model_name,
    sentiment_mode="without",
    all_dates=df.index, all_adj=df['Adj Close'],
    scaler_for_inference=scaler_for_inference,
    stock_name="TSLA"
):
    """
        + vs

    - model
    - X_train / y_train
    - X_test / y_test

```



```

- all_dates      index  df.index
- all_adj        df['Adj Close']
- scaler_for_inference  scaler  fit
- title_prefix    'NFLX'
- sentiment_mode 'with' / 'without'
"""

# Window Siz: number of lookback days
lookback = X_train.shape[1]

# Number of training samples
n_train = X_train.shape[0]

# Number of testing samples
n_test = X_test.shape[0]

# Predictions will start from the index equal to lookback (since earlier
↳ data was used for input windows)
pred_start = lookback

# Time indices for training predictions aligned with full date range
train_pts = np.arange(pred_start, pred_start + n_train)

# Time indices for test predictions aligned with full date range
test_pts = np.arange(pred_start + n_train,
                     pred_start + n_train + n_test)

# Get model predictions and true values (unscaled)
pred_close_train, _ = inv_preds(model, X_train, y_train)
pred_close_test, _ = inv_preds(model, X_test, y_test)

# Set title & label based on mode
title_suffix = f"({model_name} w/ Tweets)" if sentiment_mode == "with" else
↳ f"({model_name} w/o Tweets)"

plt.figure(figsize=(18, 7))
sns.lineplot(x=all_dates, y=all_adj, label='Actual Closing', color='black')
sns.lineplot(x=all_dates[train_pts], y=pred_close_train, label='Train
↳ Predicted', color='red')
sns.lineplot(x=all_dates[test_pts], y=pred_close_test, label='Test
↳ Predicted')

plt.title(f'{stock_name} - Closing Price vs Actual Stock {title_suffix}',
↳ fontsize=20)
plt.xlabel('Date', fontsize=18)
plt.ylabel('Price (Dollar)', fontsize=18)
plt.xticks(fontsize=14)

```

```
plt.yticks(fontsize=14)
plt.legend(fontsize=16)
plt.grid(True)
plt.tight_layout()
plt.show()
```

## Model Presetting & Building

```
[150]: # def cnn_biLSTM(input_shape, output_dim):

#     model = Sequential()

#     model.add(Conv1D(128, kernel_size=2, strides=1, padding='valid',
# ↪input_shape=input_shape))
#     model.add(MaxPooling1D(pool_size=2, strides=2))

#     model.add(Conv1D(64, kernel_size=2, strides=1, padding='valid'))
#     model.add(MaxPooling1D(pool_size=1, strides=2))

#     model.add(Bidirectional(LSTM(256, return_sequences=True)))
#     model.add(Dropout(0.2))
#     model.add(Bidirectional(LSTM(256, return_sequences=True)))
#     model.add(Dropout(0.2))

#     model.add(Dense(32, activation='relu'))
#     model.add(Dense(output_dim, activation='relu'))
#     # model.summary()
#     return model
```

```
[151]: def cnn_biLSTM(input_shape, output_dim):

    model = Sequential()

    model.add(Conv1D(128, kernel_size=2, strides=1, padding='valid',
↪input_shape=input_shape))
    model.add(MaxPooling1D(pool_size=2, strides=2))

    model.add(Conv1D(64, kernel_size=2, strides=1, padding='valid'))
    model.add(MaxPooling1D(pool_size=1, strides=2))

    model.add(Bidirectional(LSTM(256, return_sequences=True)))
    model.add(Dropout(0.2))
    model.add(Bidirectional(LSTM(256, return_sequences=True)))
    model.add(Dropout(0.2))

    model.add(Dense(32, activation='relu'))
    model.add(Dense(output_dim, activation='relu'))
    # model.summary()
```

```
return model
```

```
[152]: def directional_loss(y_true, y_pred):
        sign_true = K.sign(y_true[1:] - y_true[:-1])
        sign_pred = K.sign(y_pred[1:] - y_pred[:-1])
        correct_direction = K.cast(K.equal(sign_true, sign_pred), dtype=tf.float32)
        return 1.0 - K.mean(correct_direction)

    def integrated_loss(delta=0.1, lambda_dir=0.5):
        huber = Huber(delta=delta)
        def loss(y_true, y_pred):
            val_loss = huber(y_true, y_pred)
            dir_loss = directional_loss(y_true, y_pred)
            return (1 - lambda_dir) * val_loss + lambda_dir * dir_loss
        return loss
```

```
[153]: # Build models
cnnBiLSTM_woSent = cnn_biLSTM(
    (trainX_wo_tweet.shape[1], trainX_wo_tweet.shape[2]), trainY.shape[2]
)
cnnBiLSTM_woSent.compile(
    optimizer=Adam(learning_rate=0.0005),
    # loss=integrated_loss(delta=0.1, lambda_dir=0.16), # adjust as needed
    loss='mse', # adjust as needed
    metrics=['mae']
)

cnnBiLSTM_withSent = cnn_biLSTM(
    (trainX_with_tweet.shape[1], trainX_with_tweet.shape[2]), trainY.shape[2]
)
cnnBiLSTM_withSent.compile(
    optimizer=Adam(learning_rate=0.0005),
    # loss=integrated_loss(delta=0.1, lambda_dir=0.16), # adjust as needed
    loss='mse', # adjust as needed
    metrics=['mae']
)
```

/Users/yourth/Desktop/aaa/venv/lib/python3.11/site-packages/keras/src/layers/convolutional/base\_conv.py:107: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
[154]: early_stop = EarlyStopping(
        monitor='val_loss',
        patience=8,
```

```

        restore_best_weights=False
    )

    reduce_lr = ReduceLROnPlateau(
        monitor='val_loss',
        factor=0.7,          # ← good default
        patience=3,
        min_lr=1e-6,
        verbose=1
    )

```

## Model Fitting & Evaluating

```

[155]: # Fit models
history_cnnBiLSTM_woSent = cnnBiLSTM_woSent.fit(
    trainX_wo_tweet,
    trainY_wo_tweet,
    epochs=50,
    batch_size=64,
    validation_data=(testX_wo_tweet, testY_wo_tweet), # ← use your test split
    ↪ here
    verbose=1,
    callbacks=[early_stop, reduce_lr]
)

history_cnnBiLSTM_withSent = cnnBiLSTM_withSent.fit(
    trainX_with_tweet,
    trainY_with_tweet,
    epochs=50,
    batch_size=64,
    validation_data=(testX_with_tweet, testY_with_tweet), # ← and here
    verbose=1,
    callbacks=[early_stop, reduce_lr]
)

```

Epoch 1/50

```

2025-05-10 21:49:59.185667: E tensorflow/core/framework/node_def_util.cc:680]
NodeDef mentions attribute use_unbounded_threadpool which is not in the op
definition: Op<name=MapDataset; signature=input_dataset:variant,
other_arguments: -> handle:variant; attr=f:func;
attr=Targuments:list(type),min=0; attr=output_types:list(type),min=1;
attr=output_shapes:list(shape),min=1;
attr=use_inter_op_parallelism:bool,default=true;
attr=preserve_cardinality:bool,default=false;
attr=force_synchronous:bool,default=false; attr=metadata:string,default=""> This
may be expected if your graph generating binary is newer than this binary.
Unknown attributes will be ignored. NodeDef: {{node ParallelMapDatasetV2/_15}}
2025-05-10 21:49:59.186132: E tensorflow/core/framework/node_def_util.cc:680]

```

NodeDef mentions attribute use\_unbounded\_threadpool which is not in the op definition: Op<name=MapDataset; signature=input\_dataset:variant, other\_arguments: -> handle:variant; attr=f:func; attr=Targuments:list(type),min=0; attr=output\_types:list(type),min=1; attr=output\_shapes:list(shape),min=1; attr=use\_inter\_op\_parallelism:bool,default=true; attr=preserve\_cardinality:bool,default=false; attr=force\_synchronous:bool,default=false; attr=metadata:string,default=""> This may be expected if your graph generating binary is newer than this binary. Unknown attributes will be ignored. NodeDef: {{node ParallelMapDatasetV2/\_15}}

```

14/14                2s 26ms/step -
loss: 0.2413 - mae: 0.4284 - val_loss: 0.0749 - val_mae: 0.2474 - learning_rate:
5.0000e-04
Epoch 2/50
14/14                0s 9ms/step - loss:
0.0459 - mae: 0.1763 - val_loss: 0.0221 - val_mae: 0.1247 - learning_rate:
5.0000e-04
Epoch 3/50
14/14                0s 9ms/step - loss:
0.0158 - mae: 0.1042 - val_loss: 0.0181 - val_mae: 0.1132 - learning_rate:
5.0000e-04
Epoch 4/50
14/14                0s 9ms/step - loss:
0.0086 - mae: 0.0720 - val_loss: 0.0243 - val_mae: 0.1329 - learning_rate:
5.0000e-04
Epoch 5/50
14/14                0s 9ms/step - loss:
0.0055 - mae: 0.0585 - val_loss: 0.0178 - val_mae: 0.1139 - learning_rate:
5.0000e-04
Epoch 6/50
14/14                0s 9ms/step - loss:
0.0042 - mae: 0.0503 - val_loss: 0.0109 - val_mae: 0.0871 - learning_rate:
5.0000e-04
Epoch 7/50
14/14                0s 9ms/step - loss:
0.0041 - mae: 0.0494 - val_loss: 0.0076 - val_mae: 0.0682 - learning_rate:
5.0000e-04
Epoch 8/50
14/14                0s 8ms/step - loss:
0.0036 - mae: 0.0463 - val_loss: 0.0066 - val_mae: 0.0608 - learning_rate:
5.0000e-04
Epoch 9/50
14/14                0s 8ms/step - loss:
0.0037 - mae: 0.0471 - val_loss: 0.0063 - val_mae: 0.0586 - learning_rate:
5.0000e-04
Epoch 10/50
14/14                0s 9ms/step - loss:

```

```

0.0033 - mae: 0.0442 - val_loss: 0.0062 - val_mae: 0.0585 - learning_rate:
5.0000e-04
Epoch 11/50
14/14          0s 9ms/step - loss:
0.0034 - mae: 0.0445 - val_loss: 0.0061 - val_mae: 0.0575 - learning_rate:
5.0000e-04
Epoch 12/50
14/14          0s 9ms/step - loss:
0.0031 - mae: 0.0434 - val_loss: 0.0063 - val_mae: 0.0591 - learning_rate:
5.0000e-04
Epoch 13/50
14/14          0s 8ms/step - loss:
0.0031 - mae: 0.0434 - val_loss: 0.0062 - val_mae: 0.0584 - learning_rate:
5.0000e-04
Epoch 14/50
 9/14          0s 6ms/step - loss:
0.0030 - mae: 0.0414
Epoch 14: ReduceLROnPlateau reducing learning rate to 0.00035000001662410796.
14/14          0s 9ms/step - loss:
0.0029 - mae: 0.0414 - val_loss: 0.0060 - val_mae: 0.0577 - learning_rate:
5.0000e-04
Epoch 15/50
14/14          0s 9ms/step - loss:
0.0031 - mae: 0.0433 - val_loss: 0.0060 - val_mae: 0.0575 - learning_rate:
3.5000e-04
Epoch 16/50
14/14          0s 9ms/step - loss:
0.0031 - mae: 0.0440 - val_loss: 0.0060 - val_mae: 0.0574 - learning_rate:
3.5000e-04
Epoch 17/50
14/14          0s 9ms/step - loss:
0.0029 - mae: 0.0417 - val_loss: 0.0059 - val_mae: 0.0573 - learning_rate:
3.5000e-04
Epoch 18/50
14/14          0s 9ms/step - loss:
0.0030 - mae: 0.0420 - val_loss: 0.0059 - val_mae: 0.0575 - learning_rate:
3.5000e-04
Epoch 19/50
 9/14          0s 6ms/step - loss:
0.0032 - mae: 0.0434
Epoch 19: ReduceLROnPlateau reducing learning rate to 0.00024500001163687554.
14/14          0s 9ms/step - loss:
0.0031 - mae: 0.0430 - val_loss: 0.0059 - val_mae: 0.0573 - learning_rate:
3.5000e-04
Epoch 20/50
14/14          0s 8ms/step - loss:
0.0031 - mae: 0.0424 - val_loss: 0.0059 - val_mae: 0.0572 - learning_rate:
2.4500e-04

```

Epoch 21/50  
14/14 0s 9ms/step - loss:  
0.0028 - mae: 0.0409 - val\_loss: 0.0059 - val\_mae: 0.0571 - learning\_rate:  
2.4500e-04

Epoch 22/50  
14/14 0s 9ms/step - loss:  
0.0027 - mae: 0.0405 - val\_loss: 0.0058 - val\_mae: 0.0567 - learning\_rate:  
2.4500e-04

Epoch 23/50  
14/14 0s 9ms/step - loss:  
0.0029 - mae: 0.0412 - val\_loss: 0.0058 - val\_mae: 0.0566 - learning\_rate:  
2.4500e-04

Epoch 24/50  
14/14 0s 9ms/step - loss:  
0.0031 - mae: 0.0424 - val\_loss: 0.0058 - val\_mae: 0.0569 - learning\_rate:  
2.4500e-04

Epoch 25/50  
9/14 0s 6ms/step - loss:  
0.0033 - mae: 0.0443  
Epoch 25: ReduceLR0nPlateau reducing learning rate to 0.00017150000203400848.

14/14 0s 9ms/step - loss:  
0.0032 - mae: 0.0437 - val\_loss: 0.0058 - val\_mae: 0.0569 - learning\_rate:  
2.4500e-04

Epoch 26/50  
14/14 0s 9ms/step - loss:  
0.0028 - mae: 0.0403 - val\_loss: 0.0060 - val\_mae: 0.0576 - learning\_rate:  
1.7150e-04

Epoch 27/50  
14/14 0s 9ms/step - loss:  
0.0027 - mae: 0.0403 - val\_loss: 0.0063 - val\_mae: 0.0594 - learning\_rate:  
1.7150e-04

Epoch 28/50  
9/14 0s 6ms/step - loss:  
0.0029 - mae: 0.0413  
Epoch 28: ReduceLR0nPlateau reducing learning rate to 0.00012004999734926967.

14/14 0s 9ms/step - loss:  
0.0028 - mae: 0.0410 - val\_loss: 0.0061 - val\_mae: 0.0583 - learning\_rate:  
1.7150e-04

Epoch 29/50  
14/14 0s 13ms/step -  
loss: 0.0027 - mae: 0.0401 - val\_loss: 0.0062 - val\_mae: 0.0591 - learning\_rate:  
1.2005e-04

Epoch 30/50  
14/14 0s 9ms/step - loss:  
0.0026 - mae: 0.0393 - val\_loss: 0.0061 - val\_mae: 0.0586 - learning\_rate:  
1.2005e-04

Epoch 31/50  
9/14 0s 7ms/step - loss:

```

0.0026 - mae: 0.0386
Epoch 31: ReduceLROnPlateau reducing learning rate to 8.403499814448878e-05.
14/14          0s 9ms/step - loss:
0.0025 - mae: 0.0380 - val_loss: 0.0061 - val_mae: 0.0586 - learning_rate:
1.2005e-04
Epoch 1/50

2025-05-10 21:50:05.380287: E tensorflow/core/framework/node_def_util.cc:680]
NodeDef mentions attribute use_unbounded_threadpool which is not in the op
definition: Op<name=MapDataset; signature=input_dataset:variant,
other_arguments: -> handle:variant; attr=f:func;
attr=Targuments:list(type),min=0; attr=output_types:list(type),min=1;
attr=output_shapes:list(shape),min=1;
attr=use_inter_op_parallelism:bool,default=true;
attr=preserve_cardinality:bool,default=false;
attr=force_synchronous:bool,default=false; attr=metadata:string,default=""> This
may be expected if your graph generating binary is newer than this binary.
Unknown attributes will be ignored. NodeDef: {{node ParallelMapDatasetV2/_15}}
2025-05-10 21:50:05.380653: E tensorflow/core/framework/node_def_util.cc:680]
NodeDef mentions attribute use_unbounded_threadpool which is not in the op
definition: Op<name=MapDataset; signature=input_dataset:variant,
other_arguments: -> handle:variant; attr=f:func;
attr=Targuments:list(type),min=0; attr=output_types:list(type),min=1;
attr=output_shapes:list(shape),min=1;
attr=use_inter_op_parallelism:bool,default=true;
attr=preserve_cardinality:bool,default=false;
attr=force_synchronous:bool,default=false; attr=metadata:string,default=""> This
may be expected if your graph generating binary is newer than this binary.
Unknown attributes will be ignored. NodeDef: {{node ParallelMapDatasetV2/_15}}

14/14          2s 26ms/step -
loss: 0.1986 - mae: 0.3746 - val_loss: 0.0223 - val_mae: 0.1270 - learning_rate:
5.0000e-04
Epoch 2/50
14/14          0s 9ms/step - loss:
0.0262 - mae: 0.1320 - val_loss: 0.0524 - val_mae: 0.1963 - learning_rate:
5.0000e-04
Epoch 3/50
14/14          0s 9ms/step - loss:
0.0118 - mae: 0.0881 - val_loss: 0.0615 - val_mae: 0.2217 - learning_rate:
5.0000e-04
Epoch 4/50
 9/14          0s 6ms/step - loss:
0.0075 - mae: 0.0664
Epoch 4: ReduceLROnPlateau reducing learning rate to 0.00035000001662410796.
14/14          0s 9ms/step - loss:
0.0070 - mae: 0.0640 - val_loss: 0.0267 - val_mae: 0.1430 - learning_rate:
5.0000e-04
Epoch 5/50

```



```

14/14          0s 9ms/step - loss:
0.0048 - mae: 0.0550 - val_loss: 0.0142 - val_mae: 0.1010 - learning_rate:
3.5000e-04
Epoch 6/50
14/14          0s 9ms/step - loss:
0.0038 - mae: 0.0477 - val_loss: 0.0075 - val_mae: 0.0674 - learning_rate:
3.5000e-04
Epoch 7/50
14/14          0s 8ms/step - loss:
0.0037 - mae: 0.0470 - val_loss: 0.0065 - val_mae: 0.0599 - learning_rate:
3.5000e-04
Epoch 8/50
14/14          0s 9ms/step - loss:
0.0035 - mae: 0.0461 - val_loss: 0.0063 - val_mae: 0.0587 - learning_rate:
3.5000e-04
Epoch 9/50
14/14          0s 9ms/step - loss:
0.0034 - mae: 0.0446 - val_loss: 0.0062 - val_mae: 0.0583 - learning_rate:
3.5000e-04
Epoch 10/50
14/14          0s 9ms/step - loss:
0.0032 - mae: 0.0432 - val_loss: 0.0061 - val_mae: 0.0582 - learning_rate:
3.5000e-04
Epoch 11/50
14/14          0s 10ms/step -
loss: 0.0030 - mae: 0.0418 - val_loss: 0.0062 - val_mae: 0.0587 - learning_rate:
3.5000e-04
Epoch 12/50
 9/14          0s 7ms/step - loss:
0.0029 - mae: 0.0407
Epoch 12: ReduceLROnPlateau reducing learning rate to 0.00024500001163687554.
14/14          0s 9ms/step - loss:
0.0029 - mae: 0.0409 - val_loss: 0.0061 - val_mae: 0.0582 - learning_rate:
3.5000e-04
Epoch 13/50
14/14          0s 10ms/step -
loss: 0.0030 - mae: 0.0421 - val_loss: 0.0061 - val_mae: 0.0584 - learning_rate:
2.4500e-04
Epoch 14/50
14/14          0s 9ms/step - loss:
0.0030 - mae: 0.0425 - val_loss: 0.0061 - val_mae: 0.0587 - learning_rate:
2.4500e-04
Epoch 15/50
14/14          0s 9ms/step - loss:
0.0031 - mae: 0.0433 - val_loss: 0.0061 - val_mae: 0.0583 - learning_rate:
2.4500e-04
Epoch 16/50
14/14          0s 9ms/step - loss:

```

```

0.0029 - mae: 0.0412 - val_loss: 0.0060 - val_mae: 0.0576 - learning_rate:
2.4500e-04
Epoch 17/50
14/14          0s 9ms/step - loss:
0.0029 - mae: 0.0423 - val_loss: 0.0062 - val_mae: 0.0588 - learning_rate:
2.4500e-04
Epoch 18/50
14/14          0s 9ms/step - loss:
0.0031 - mae: 0.0422 - val_loss: 0.0059 - val_mae: 0.0575 - learning_rate:
2.4500e-04
Epoch 19/50
 9/14          0s 7ms/step - loss:
0.0028 - mae: 0.0409
Epoch 19: ReduceLROnPlateau reducing learning rate to 0.00017150000203400848.
14/14          0s 9ms/step - loss:
0.0027 - mae: 0.0403 - val_loss: 0.0060 - val_mae: 0.0578 - learning_rate:
2.4500e-04
Epoch 20/50
14/14          0s 9ms/step - loss:
0.0027 - mae: 0.0402 - val_loss: 0.0059 - val_mae: 0.0575 - learning_rate:
1.7150e-04
Epoch 21/50
14/14          0s 10ms/step -
loss: 0.0029 - mae: 0.0408 - val_loss: 0.0058 - val_mae: 0.0568 - learning_rate:
1.7150e-04
Epoch 22/50
14/14          0s 9ms/step - loss:
0.0028 - mae: 0.0404 - val_loss: 0.0059 - val_mae: 0.0575 - learning_rate:
1.7150e-04
Epoch 23/50
14/14          0s 10ms/step -
loss: 0.0031 - mae: 0.0427 - val_loss: 0.0060 - val_mae: 0.0583 - learning_rate:
1.7150e-04
Epoch 24/50
 9/14          0s 7ms/step - loss:
0.0032 - mae: 0.0441
Epoch 24: ReduceLROnPlateau reducing learning rate to 0.00012004999734926967.
14/14          0s 10ms/step -
loss: 0.0031 - mae: 0.0434 - val_loss: 0.0059 - val_mae: 0.0572 - learning_rate:
1.7150e-04
Epoch 25/50
14/14          0s 10ms/step -
loss: 0.0030 - mae: 0.0422 - val_loss: 0.0059 - val_mae: 0.0573 - learning_rate:
1.2005e-04
Epoch 26/50
14/14          0s 9ms/step - loss:
0.0030 - mae: 0.0419 - val_loss: 0.0062 - val_mae: 0.0591 - learning_rate:
1.2005e-04

```

```

Epoch 27/50
  9/14          0s 6ms/step - loss:
0.0027 - mae: 0.0399
Epoch 27: ReduceLROnPlateau reducing learning rate to 8.403499814448878e-05.
 14/14          0s 9ms/step - loss:
0.0027 - mae: 0.0396 - val_loss: 0.0063 - val_mae: 0.0595 - learning_rate:
1.2005e-04
Epoch 28/50
 14/14          0s 10ms/step -
loss: 0.0028 - mae: 0.0410 - val_loss: 0.0063 - val_mae: 0.0598 - learning_rate:
8.4035e-05
Epoch 29/50
 14/14          0s 10ms/step -
loss: 0.0027 - mae: 0.0402 - val_loss: 0.0059 - val_mae: 0.0574 - learning_rate:
8.4035e-05

```

```

[156]: plot_metrics(cnnBiLSTM_woSent, history_cnnBiLSTM_woSent,
                    trainX_wo_tweet, trainY_wo_tweet,
                    testX_wo_tweet, testY_wo_tweet,
                    "without")

plot_metrics(cnnBiLSTM_withSent, history_cnnBiLSTM_withSent,
            trainX_with_tweet, trainY_with_tweet,
            testX_with_tweet, testY_with_tweet,
            "with")

```

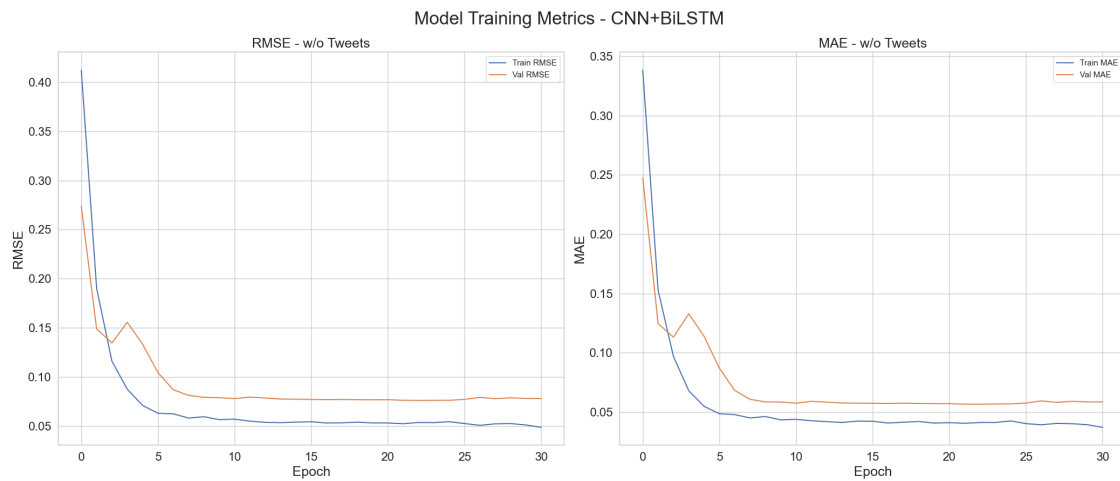
```

2025-05-10 21:50:11.310491: E tensorflow/core/framework/node_def_util.cc:680]
NodeDef mentions attribute use_unbounded_threadpool which is not in the op
definition: Op<name=MapDataset; signature=input_dataset:variant,
other_arguments: -> handle:variant; attr=f:func;
attr=Targuments:list(type),min=0; attr=output_types:list(type),min=1;
attr=output_shapes:list(shape),min=1;
attr=use_inter_op_parallelism:bool,default=true;
attr=preserve_cardinality:bool,default=false;
attr=force_synchronous:bool,default=false; attr=metadata:string,default=""> This
may be expected if your graph generating binary is newer than this binary.
Unknown attributes will be ignored. NodeDef: {{node ParallelMapDatasetV2/_14}}
2025-05-10 21:50:11.310805: E tensorflow/core/framework/node_def_util.cc:680]
NodeDef mentions attribute use_unbounded_threadpool which is not in the op
definition: Op<name=MapDataset; signature=input_dataset:variant,
other_arguments: -> handle:variant; attr=f:func;
attr=Targuments:list(type),min=0; attr=output_types:list(type),min=1;
attr=output_shapes:list(shape),min=1;
attr=use_inter_op_parallelism:bool,default=true;
attr=preserve_cardinality:bool,default=false;
attr=force_synchronous:bool,default=false; attr=metadata:string,default=""> This
may be expected if your graph generating binary is newer than this binary.
Unknown attributes will be ignored. NodeDef: {{node ParallelMapDatasetV2/_14}}

```


w/o Tweets RMSE: Train = 0.0465, Test = 0.0780

w/o Tweets MAE : Train = 0.0361, Test = 0.0586



w/ Tweets RMSE: Train = 0.0454, Test = 0.0770

w/ Tweets MAE : Train = 0.0350, Test = 0.0574



```
[157]: plot_pred_vs_actual(  
    model=cnnBiLSTM_woSent,  
    X_train=trainX_wo_tweet, y_train=trainY_wo_tweet,  
    X_test=testX_wo_tweet, y_test=testY_wo_tweet,  
    model_name="CNN+BiLSTM",  
    sentiment_mode="without"  
)  
  
plot_pred_vs_actual(  

```

```

model=cnnBiLSTM_withSent,
X_train=trainX_with_tweet, y_train=trainY_with_tweet,
X_test=testX_with_tweet, y_test=testY_with_tweet,
model_name="CNN+BiLSTM",
sentiment_mode="with"
)

```



## 2.0.3 Transformer

### Pre-load Plotting Functions

```

[93]: def plot_transformer_metrics(model, history, X_train, y_train, X_val, y_val,
    ↪ sentiment_mode):
    """
    Plot training and validation RMSE/MAE metrics and loss curves for
    ↪ Transformer models.
    """
    history_data = history.history

```

```

# 1. Predictions
y_train_pred = model.predict(X_train, verbose=0).squeeze()
y_val_pred    = model.predict(X_val, verbose=0).squeeze()

y_train_true = y_train.squeeze()
y_val_true   = y_val.squeeze()

# 2. Metrics
train_rmse = np.sqrt(mean_squared_error(y_train_true, y_train_pred))
val_rmse    = np.sqrt(mean_squared_error(y_val_true, y_val_pred))
train_mae   = mean_absolute_error(y_train_true, y_train_pred)
val_mae     = mean_absolute_error(y_val_true, y_val_pred)

title_suffix = " w/ Tweets" if sentiment_mode == "with" else "w/o Tweets"

# 3. Print metrics
print(f" [Transformer{title_suffix}] RMSE: Train = {train_rmse:.4f}, Test_
↪ {val_rmse:.4f}")
print(f" [Transformer{title_suffix}] MAE : Train = {train_mae:.4f}, Test =_
↪ {val_mae:.4f}")

# 4. Plot
plt.figure(figsize=(21, 9))

# ---- RMSE subplot
plt.subplot(1, 2, 1)
plt.plot(np.sqrt(history_data['loss']), label='Train RMSE')
plt.plot(np.sqrt(history_data['val_loss']), label='Val RMSE')
plt.title(f'RMSE - Transformer {title_suffix}', fontsize=18)
plt.xlabel('Epoch', fontsize=18)
plt.ylabel('RMSE', fontsize=18)
plt.xticks(fontsize=16)
plt.yticks(fontsize=16)
plt.legend()

# ---- MAE subplot
plt.subplot(1, 2, 2)
if 'mae' in history_data:
    plt.plot(history_data['mae'], label='Train MAE')
    plt.plot(history_data['val_mae'], label='Val MAE')
else:
    plt.plot([train_mae] * len(history_data['loss']), label='Train MAE_
↪ (flat)')
    plt.plot([val_mae] * len(history_data['val_loss']), label='Val MAE_
↪ (flat)')
plt.title(f'MAE - Transformer {title_suffix}', fontsize=18)

```

```

plt.xlabel('Epoch', fontsize=18)
plt.ylabel('MAE', fontsize=18)
plt.xticks(fontsize=16)
plt.yticks(fontsize=16)
plt.legend()

plt.suptitle('Model Training Metrics - Transformer', fontsize=24)
plt.tight_layout()
plt.show()

```

```

[94]: def plot_transformer_predictions(dates, true_prices, train_preds, test_preds,
    ↪ train_size, title):
    """
    Plot true prices vs. transformer predictions using seaborn styling.

    Args:
        dates: full date index
        true_prices: full true values (1D)
        train_preds: model predictions on training data
        test_preds: model predictions on test data
        train_size: number of training samples (for separation)
        title: chart title
    """
    # Set seaborn style
    sns.set_theme(style="whitegrid")

    # Create figure
    plt.figure(figsize=(18, 7))

    # Plot actual prices
    sns.lineplot(x=dates, y=true_prices, color='black', label='Actual Closing')

    # Plot transformer predictions
    sns.lineplot(x=dates[:train_size], y=train_preds, color='red', label='Train_
    ↪ Predicted')

    # Plot test predictions with markers
    sns.lineplot(x=dates[train_size:], y=test_preds,
        label='Test Predicted')

    # Add shaded region to distinguish train/test split
    plt.axvline(x=dates[train_size-1], color='gray', linestyle='--', alpha=0.7)

    # Styling
    plt.title(title, fontsize=20)
    plt.xlabel('Date (Jan 2020 - July 2022)', fontsize=18)
    plt.ylabel('Price (Dollar)', fontsize=18)

```

```

plt.xticks(fontsize=14)
plt.yticks(fontsize=14)
plt.legend(fontsize=16)
plt.grid(True)
plt.tight_layout()

return plt

```

## Model Presetting & Building

```

[95]: class PositionalEncoding(tf.keras.layers.Layer):
    def __init__(self, sequence_len, d_model):
        super().__init__()
        self.pos_encoding = self.positional_encoding(sequence_len, d_model)

    def get_config(self):
        return {"sequence_len": self.pos_encoding.shape[0], "d_model": self.
↪pos_encoding.shape[1]}

    def positional_encoding(self, position, d_model):
        angle_rads = self.get_angles(np.arange(position)[: , np.newaxis],
                                     np.arange(d_model)[np.newaxis, :],
                                     d_model)

        angle_rads[:, 0::2] = np.sin(angle_rads[:, 0::2])
        angle_rads[:, 1::2] = np.cos(angle_rads[:, 1::2])
        return tf.cast(angle_rads[np.newaxis, ...], dtype=tf.float32)

    def get_angles(self, pos, i, d_model):
        angle_rates = 1 / np.power(10000, (2 * (i // 2)) / np.float32(d_model))
        return pos * angle_rates

    def call(self, x):
        return x + self.pos_encoding[:, :tf.shape(x)[1], :]

```

```

[96]: # ===== 1. Transformer Encoder =====
def transformer_encoder(inputs, head_size, num_heads, ff_dim, dropout):
    """
    Builds a single Transformer encoder block.
    """

    # Multi-head self-attention
    attention_output = MultiHeadAttention(num_heads=num_heads,
↪key_dim=head_size, dropout=dropout)(inputs, inputs)
    attention_output = Add()([inputs, attention_output])
    attention_output = LayerNormalization()(attention_output)

    # Feed-forward network
    ffn_output = Dense(ff_dim, activation='relu')(attention_output)

```



```

    # ffn_output = Dropout(dropout)(ffn_output) # <--- Dropout after first FFN
    ↪ layer
    ffn_output = Dense(inputs.shape[-1])(ffn_output)
    ffn_output = Add()([attention_output, ffn_output])
    output = LayerNormalization()(ffn_output)

    return output

```

```

[97]: def build_transformer_model(input_shape, head_size=64, num_heads=4, ff_dim=128,
    ↪ num_layers=2, dropout=0.1):
    inputs = Input(shape=input_shape)
    x = PositionalEncoding(input_shape[0], input_shape[1])(inputs) # add
    ↪ positional encoding

    for _ in range(num_layers):
        x = transformer_encoder(x, head_size, num_heads, ff_dim, dropout)

    x = GlobalAveragePooling1D()(x)
    outputs = Dense(1)(x)
    return Model(inputs, outputs)

```

```

[98]: def directional_loss(y_true, y_pred):
    sign_true = K.sign(y_true[1:] - y_true[:-1])
    sign_pred = K.sign(y_pred[1:] - y_pred[:-1])
    correct_direction = K.cast(K.equal(sign_true, sign_pred), dtype=tf.float32)
    return 1.0 - K.mean(correct_direction)

def integrated_loss(delta=0.01, lambda_dir=0.5):
    huber = Huber(delta=delta)

    def loss(y_true, y_pred):
        huber_loss = huber(y_true, y_pred)
        dir_loss = directional_loss(y_true, y_pred)
        return huber_loss + lambda_dir * dir_loss

    return loss

```

```

[99]: # For Transformer *without* sentiment (cleaner input, stop sooner)
early_stop_wo = EarlyStopping(
    monitor='val_loss',
    patience=10,
    restore_best_weights=True
)

# For Transformer *with* sentiment (noisier input, allow more time)
early_stop_with = EarlyStopping(
    monitor='val_loss',

```

```

        patience=10,
        restore_best_weights=True
    )

    reduce_lr = ReduceLROnPlateau(
        monitor='val_loss',
        factor=0.75,          # ← good default
        patience=3,
        min_lr=1e-5,
        verbose=1
    )

```

## Model Training & Evaluating

```

[100]: # ===== Train model without sentiment =====
model_wo_sent = build_transformer_model((trainX_wo_tweet.shape[1],
    ↪trainX_wo_tweet.shape[2]))
model_wo_sent.compile(
    optimizer=Adam(0.001),
    # loss=integrated_loss(),
    loss=Huber(0.01),
    # loss='mse',
    metrics=['mae']
)
history_wo_sent = model_wo_sent.fit( # save history here
    trainX_wo_tweet, trainY_wo_tweet,
    validation_data=(testX_wo_tweet, testY_wo_tweet),
    epochs=50, batch_size=64, verbose=1,
    callbacks=[early_stop_wo]
)

# ===== Train model with sentiment =====
model_with_sent = build_transformer_model((trainX_with_tweet.shape[1],
    ↪trainX_with_tweet.shape[2]))
model_with_sent.compile(
    optimizer=Adam(0.001),
    # loss=integrated_loss(),
    loss=Huber(0.01),
    # loss='mse',
    metrics=['mae']
)
history_with_sent = model_with_sent.fit( # save history here
    trainX_with_tweet, trainY_with_tweet,
    validation_data=(testX_with_tweet, testY_with_tweet),
    epochs=50, batch_size=64, verbose=1,
    callbacks=[early_stop_with]
)

```

Epoch 1/50  
14/14 2s 19ms/step -  
loss: 0.0032 - mae: 0.3247 - val\_loss: 0.0049 - val\_mae: 0.4907  
Epoch 2/50  
14/14 0s 6ms/step - loss:  
0.0011 - mae: 0.1165 - val\_loss: 0.0038 - val\_mae: 0.3814  
Epoch 3/50  
14/14 0s 6ms/step - loss:  
6.6015e-04 - mae: 0.0709 - val\_loss: 0.0030 - val\_mae: 0.3041  
Epoch 4/50  
14/14 0s 6ms/step - loss:  
5.4238e-04 - mae: 0.0590 - val\_loss: 0.0030 - val\_mae: 0.3053  
Epoch 5/50  
14/14 0s 7ms/step - loss:  
4.4475e-04 - mae: 0.0492 - val\_loss: 0.0023 - val\_mae: 0.2378  
Epoch 6/50  
14/14 0s 6ms/step - loss:  
3.6234e-04 - mae: 0.0410 - val\_loss: 0.0020 - val\_mae: 0.2000  
Epoch 7/50  
14/14 0s 6ms/step - loss:  
3.4177e-04 - mae: 0.0389 - val\_loss: 0.0018 - val\_mae: 0.1817  
Epoch 8/50  
14/14 0s 6ms/step - loss:  
3.4357e-04 - mae: 0.0391 - val\_loss: 0.0017 - val\_mae: 0.1751  
Epoch 9/50  
14/14 0s 6ms/step - loss:  
2.7825e-04 - mae: 0.0324 - val\_loss: 0.0015 - val\_mae: 0.1581  
Epoch 10/50  
14/14 0s 6ms/step - loss:  
2.9995e-04 - mae: 0.0346 - val\_loss: 0.0017 - val\_mae: 0.1704  
Epoch 11/50  
14/14 0s 6ms/step - loss:  
3.0187e-04 - mae: 0.0349 - val\_loss: 0.0015 - val\_mae: 0.1548  
Epoch 12/50  
14/14 0s 6ms/step - loss:  
2.9833e-04 - mae: 0.0345 - val\_loss: 0.0018 - val\_mae: 0.1818  
Epoch 13/50  
14/14 0s 6ms/step - loss:  
3.3264e-04 - mae: 0.0379 - val\_loss: 0.0017 - val\_mae: 0.1767  
Epoch 14/50  
14/14 0s 6ms/step - loss:  
2.6231e-04 - mae: 0.0308 - val\_loss: 0.0016 - val\_mae: 0.1627  
Epoch 15/50  
14/14 0s 6ms/step - loss:  
3.2832e-04 - mae: 0.0376 - val\_loss: 0.0016 - val\_mae: 0.1661  
Epoch 16/50  
14/14 0s 7ms/step - loss:  
2.5449e-04 - mae: 0.0301 - val\_loss: 0.0016 - val\_mae: 0.1694

```

Epoch 17/50
14/14          0s 6ms/step - loss:
2.7043e-04 - mae: 0.0317 - val_loss: 0.0015 - val_mae: 0.1599
Epoch 18/50
14/14          0s 6ms/step - loss:
3.1782e-04 - mae: 0.0365 - val_loss: 0.0013 - val_mae: 0.1374
Epoch 19/50
14/14          0s 6ms/step - loss:
3.7832e-04 - mae: 0.0426 - val_loss: 0.0013 - val_mae: 0.1358
Epoch 20/50
14/14          0s 6ms/step - loss:
2.9102e-04 - mae: 0.0337 - val_loss: 0.0016 - val_mae: 0.1673
Epoch 21/50
14/14          0s 6ms/step - loss:
2.6692e-04 - mae: 0.0313 - val_loss: 0.0017 - val_mae: 0.1762
Epoch 22/50
14/14          0s 6ms/step - loss:
2.8483e-04 - mae: 0.0331 - val_loss: 0.0016 - val_mae: 0.1664
Epoch 23/50
14/14          0s 6ms/step - loss:
2.4354e-04 - mae: 0.0290 - val_loss: 0.0017 - val_mae: 0.1749
Epoch 24/50
14/14          0s 6ms/step - loss:
2.1451e-04 - mae: 0.0261 - val_loss: 0.0017 - val_mae: 0.1715
Epoch 25/50
14/14          0s 6ms/step - loss:
2.3525e-04 - mae: 0.0282 - val_loss: 0.0018 - val_mae: 0.1808
Epoch 26/50
14/14          0s 6ms/step - loss:
2.2279e-04 - mae: 0.0269 - val_loss: 0.0017 - val_mae: 0.1799
Epoch 27/50
14/14          0s 5ms/step - loss:
2.1414e-04 - mae: 0.0259 - val_loss: 0.0017 - val_mae: 0.1731
Epoch 28/50
14/14          0s 6ms/step - loss:
2.0664e-04 - mae: 0.0252 - val_loss: 0.0020 - val_mae: 0.2022
Epoch 29/50
14/14          0s 5ms/step - loss:
2.5003e-04 - mae: 0.0297 - val_loss: 0.0019 - val_mae: 0.1957
Epoch 1/50

```

```

2025-05-10 13:51:53.273245: E tensorflow/core/framework/node_def_util.cc:680]
NodeDef mentions attribute use_unbounded_threadpool which is not in the op
definition: Op<name=MapDataset; signature=input_dataset:variant,
other_arguments: -> handle:variant; attr=f:func;
attr=Targuments:list(type),min=0; attr=output_types:list(type),min=1;
attr=output_shapes:list(shape),min=1;
attr=use_inter_op_parallelism:bool,default=true;

```

```

attr=preserve_cardinality:bool,default=false;
attr=force_synchronous:bool,default=false; attr=metadata:string,default=""> This
may be expected if your graph generating binary is newer than this binary.
Unknown attributes will be ignored. NodeDef: {{node ParallelMapDatasetV2/_15}}
2025-05-10 13:51:53.273545: E tensorflow/core/framework/node_def_util.cc:680]
NodeDef mentions attribute use_unbounded_threadpool which is not in the op
definition: Op<name=MapDataset; signature=input_dataset:variant,
other_arguments: -> handle:variant; attr=f:func;
attr=Targuments:list(type),min=0; attr=output_types:list(type),min=1;
attr=output_shapes:list(shape),min=1;
attr=use_inter_op_parallelism:bool,default=true;
attr=preserve_cardinality:bool,default=false;
attr=force_synchronous:bool,default=false; attr=metadata:string,default=""> This
may be expected if your graph generating binary is newer than this binary.
Unknown attributes will be ignored. NodeDef: {{node ParallelMapDatasetV2/_15}}

```

```

14/14          2s 18ms/step -
loss: 0.0030 - mae: 0.3078 - val_loss: 0.0059 - val_mae: 0.5982
Epoch 2/50
14/14          0s 7ms/step - loss:
0.0011 - mae: 0.1116 - val_loss: 0.0015 - val_mae: 0.1580
Epoch 3/50
14/14          0s 7ms/step - loss:
6.0356e-04 - mae: 0.0652 - val_loss: 0.0018 - val_mae: 0.1837
Epoch 4/50
14/14          0s 7ms/step - loss:
4.0895e-04 - mae: 0.0457 - val_loss: 0.0012 - val_mae: 0.1277
Epoch 5/50
14/14          0s 7ms/step - loss:
3.1472e-04 - mae: 0.0362 - val_loss: 8.5065e-04 - val_mae: 0.0900
Epoch 6/50
14/14          0s 7ms/step - loss:
2.8180e-04 - mae: 0.0328 - val_loss: 7.5442e-04 - val_mae: 0.0804
Epoch 7/50
14/14          0s 6ms/step - loss:
2.4849e-04 - mae: 0.0295 - val_loss: 8.4502e-04 - val_mae: 0.0893
Epoch 8/50
14/14          0s 6ms/step - loss:
2.8955e-04 - mae: 0.0337 - val_loss: 0.0012 - val_mae: 0.1210
Epoch 9/50
14/14          0s 6ms/step - loss:
2.3295e-04 - mae: 0.0279 - val_loss: 8.5472e-04 - val_mae: 0.0904
Epoch 10/50
14/14          0s 6ms/step - loss:
2.6701e-04 - mae: 0.0314 - val_loss: 0.0014 - val_mae: 0.1423
Epoch 11/50
14/14          0s 6ms/step - loss:
2.3613e-04 - mae: 0.0282 - val_loss: 0.0012 - val_mae: 0.1222

```

```

Epoch 12/50
14/14          0s 6ms/step - loss:
2.8167e-04 - mae: 0.0329 - val_loss: 0.0011 - val_mae: 0.1191
Epoch 13/50
14/14          0s 6ms/step - loss:
2.0412e-04 - mae: 0.0250 - val_loss: 0.0011 - val_mae: 0.1144
Epoch 14/50
14/14          0s 6ms/step - loss:
3.3401e-04 - mae: 0.0381 - val_loss: 8.5530e-04 - val_mae: 0.0903
Epoch 15/50
14/14          0s 6ms/step - loss:
3.3361e-04 - mae: 0.0381 - val_loss: 9.6195e-04 - val_mae: 0.1011
Epoch 16/50
14/14          0s 6ms/step - loss:
2.9512e-04 - mae: 0.0342 - val_loss: 9.0472e-04 - val_mae: 0.0954

```

```


[101]: plot_transformer_metrics(
        model=model_wo_sent,
        history=history_wo_sent,
        X_train=trainX_wo_tweet, y_train=trainY_wo_tweet,
        X_val=testX_wo_tweet, y_val=testY_wo_tweet,
        sentiment_mode="without"
    )

plot_transformer_metrics(
    model=model_with_sent,
    history=history_with_sent,
    X_train=trainX_with_tweet, y_train=trainY_with_tweet,
    X_val=testX_with_tweet, y_val=testY_with_tweet,
    sentiment_mode="with"
)

```

[Transformerw/o Tweets] RMSE: Train = 0.0376, Test = 0.1780

[Transformerw/o Tweets] MAE : Train = 0.0295, Test = 0.1358



[Transformer w/ Tweets] RMSE: Train = 0.0345, Test = 0.0944

[Transformer w/ Tweets] MAE : Train = 0.0259, Test = 0.0804



```
[102]: # ===== 1. Get Raw Predictions =====

# Model without sentiment
train_preds_raw_wo_sent = model_wo_sent.predict(trainX_wo_tweet).reshape(-1, 1)
test_preds_raw_wo_sent  = model_wo_sent.predict(testX_wo_tweet).reshape(-1, 1)

# Model with sentiment
train_preds_raw_with_sent = model_with_sent.predict(trainX_with_tweet).
    ↳ reshape(-1, 1)
test_preds_raw_with_sent  = model_with_sent.predict(testX_with_tweet).
    ↳ reshape(-1, 1)

# ===== 2. Inverse Transform Predictions =====

def safe_inverse(preds_scaled, full_scaler, target_index=0):
    """
    Inverse-transform predictions scaled with a full-feature MinMaxScaler.
    Assumes predictions correspond to feature at `target_index`.
    """
    pad_width = full_scaler.n_features_in_ - 1
    padded = np.concatenate([preds_scaled, np.zeros((preds_scaled.shape[0],
    ↳ pad_width))], axis=1)
    return full_scaler.inverse_transform(padded[:, target_index]).reshape(-1, 1)

# Use the safe inverse
```

```

inv_train_preds_wo_sent = safe_inverse(train_preds_raw_wo_sent, scaler,
    ↪target_index=0)
inv_test_preds_wo_sent = safe_inverse(test_preds_raw_wo_sent, scaler,
    ↪target_index=0)

inv_train_preds_with_sent = safe_inverse(train_preds_raw_with_sent, scaler,
    ↪target_index=0)
inv_test_preds_with_sent = safe_inverse(test_preds_raw_with_sent, scaler,
    ↪target_index=0)

# ===== 3. Prepare Date Range and Ground Truth =====

# For model without sentiment
train_size_wo = len(inv_train_preds_wo_sent)
total_preds_wo = train_size_wo + len(inv_test_preds_wo_sent)
dates_wo = df_filtered.index[n_past:n_past + total_preds_wo]
true_adj_close_wo = df_filtered['Adj Close'].values[n_past:n_past +
    ↪total_preds_wo]

# For model with sentiment
train_size_with = len(inv_train_preds_with_sent)
total_preds_with = train_size_with + len(inv_test_preds_with_sent)
dates_with = df_filtered.index[n_past:n_past + total_preds_with]
true_adj_close_with = df_filtered['Adj Close'].values[n_past:n_past +
    ↪total_preds_with]

# ===== 4. Plot Predictions =====

# Plot: WITHOUT Sentiment
plot_transformer_predictions(
    dates=dates_wo,
    true_prices=true_adj_close_wo,
    train_preds=inv_train_preds_wo_sent.flatten(),
    test_preds=inv_test_preds_wo_sent.flatten(),
    train_size=train_size_wo,
    title="TSLA - Closing Price vs Actual Stock (Transformer w/o Tweets)"
).show()


# Plot: WITH Sentiment
plot_transformer_predictions(
    dates=dates_with,
    true_prices=true_adj_close_with,
    train_preds=inv_train_preds_with_sent.flatten(),
    test_preds=inv_test_preds_with_sent.flatten(),
    train_size=train_size_with,
    title="TSLA - Closing Price vs Actual Stock (Transformer w/ Tweets)"
)

```



```
).show()
```

```
28/28      0s 2ms/step
12/12      0s 2ms/step
28/28      0s 2ms/step
12/12      0s 2ms/step
```



```
[103]: def evaluate_prediction_volatility(pred1, pred2, true=None, name1="Model A",  
      ↪name2="Model B"):  
      """  
  
      - pred1, pred2:      1D numpy array list  
      - true:  
      """
```

```

def metrics(arr):
    arr = np.array(arr)
    first_diff = np.diff(arr)
    std = np.std(first_diff)
    mad = np.mean(np.abs(first_diff))
    smooth = np.mean(np.abs(arr[2:] - 2 * arr[1:-1] + arr[:-2]))
    return std, mad, smooth

std1, mad1, smooth1 = metrics(pred1)
std2, mad2, smooth2 = metrics(pred2)

print(f" {name1}      :")
print(f"      ·      std(diff)          = {std1:.6f}")
print(f"      ·      mean(abs) = {mad1:.6f}")
print(f"      ·      smoothness          = {smooth1:.6f}")
print()

print(f" {name2}      :")
print(f"      ·      std(diff)          = {std2:.6f}")
print(f"      ·      mean(abs)          = {mad2:.6f}")
print(f"      ·      smoothness          = {smooth2:.6f}")
print()

if true is not None:
    std_t, mad_t, smooth_t = metrics(true)
    print(f"      :")
    print(f"      · std(diff)          = {std_t:.6f}")
    print(f"      · mean(abs)          = {mad_t:.6f}")
    print(f"      · smoothness          = {smooth_t:.6f}")
    print()

```

```

[104]: pred1 = np.concatenate([inv_train_preds_wo_sent, inv_test_preds_wo_sent]).
        ↪flatten()
pred2 = np.concatenate([inv_train_preds_with_sent, inv_test_preds_with_sent]).
        ↪flatten()

evaluate_prediction_volatility(
    pred1=pred1,
    pred2=pred2,
    true=true_adj_close_wo,
    name1="Transformer w/o sent",
    name2="Transformer w/ sent"
)

```

```

Transformer w/o sent      :
·      std(diff)          = 1.677407
·      mean(abs) = 1.160032
·      smoothness          = 1.674059

```

```
Transformer w/ sent      :
·   std(diff)            = 1.314947
·   mean(abs)            = 1.017639
·   smoothness           = 1.420668

:
·   std(diff)            = 1.129569
·   mean(abs)            = 0.742417
·   smoothness           = 1.118966
```

```
[105]: pred1 = np.concatenate([inv_train_preds_wo_sent, inv_test_preds_wo_sent]).
        ↪flatten()
pred2 = np.concatenate([inv_train_preds_with_sent, inv_test_preds_with_sent]).
        ↪flatten()

evaluate_prediction_volatility(
    pred1=pred1,
    pred2=pred2,
    true=true_adj_close_wo,
    name1="Transformer w/o sent",
    name2="Transformer w/ sent"
)
```

```
Transformer w/o sent      :
·   std(diff)            = 1.677407
·   mean(abs)            = 1.160032
·   smoothness           = 1.674059
```

```
Transformer w/ sent      :
·   std(diff)            = 1.314947
·   mean(abs)            = 1.017639
·   smoothness           = 1.420668

:
·   std(diff)            = 1.129569
·   mean(abs)            = 0.742417
·   smoothness           = 1.118966
```

```
[106]: pred_close_train_wo, _ = inv_preds(cnnBiLSTM_woSent, trainX_wo_tweet,
        ↪trainY_wo_tweet)
pred_close_test_wo, _ = inv_preds(cnnBiLSTM_woSent, testX_wo_tweet,
        ↪testY_wo_tweet)

pred_close_train_w, _ = inv_preds(cnnBiLSTM_withSent, trainX_with_tweet,
        ↪trainY_with_tweet)
```

```

pred_close_test_w, _ = inv_preds(cnnBiLSTM_withSent, testX_with_tweet,
    ↪testY_with_tweet)

pred3 = np.concatenate([pred_close_train_wo, pred_close_test_wo]).flatten()
pred4 = np.concatenate([pred_close_train_w, pred_close_test_w]).flatten()

evaluate_prediction_volatility(
    pred1=pred3,
    pred2=pred4,
    true=true_adj_close_wo,
    name1="cnnBiLSTM w/o sent",
    name2="cnnBiLSTM w/ sent"
)

```

```

cnnBiLSTM w/o sent      :
·      std(diff)        = 0.689062
·      mean(abs)        = 0.514961
·      smoothness       = 0.676101

```

```

cnnBiLSTM w/ sent      :
·      std(diff)        = 1.031721
·      mean(abs)        = 0.802738
·      smoothness       = 1.240502

```

```

:
·      std(diff)        = 1.129569
·      mean(abs)        = 0.742417
·      smoothness       = 1.118966

```

```


[107]: from statsmodels.graphics.tsaplots import plot_acf
import matplotlib.pyplot as plt
import numpy as np

def plot_acf_diff(series, label):
    diff = np.diff(series)
    plt.figure(figsize=(10, 4))
    plot_acf(diff, lags=40, title=f"ACF of First Difference - {label}")
    plt.grid(True)
    plt.show()


plot_acf_diff(pred1, "Transformer w/o Sentiment")
plot_acf_diff(pred2, "Transformer w/ Sentiment")
plot_acf_diff(true_adj_close_wo, "True Closing Price")

```


<Figure size 1000x400 with 0 Axes>



<Figure size 1000x400 with 0 Axes>



<Figure size 1000x400 with 0 Axes>



```
[108]: from scipy.fft import fft, fftfreq
import matplotlib.pyplot as plt

def plot_frequency_spectrum(series, label, sample_rate=1):
    diff = np.diff(series)
    N = len(diff)
    yf = np.abs(fft(diff))
    xf = fftfreq(N, d=sample_rate)[:N // 2]

    plt.figure(figsize=(10, 4))
    plt.plot(xf, yf[:N // 2])
    plt.title(f"Frequency Spectrum of 1st Difference - {label}")
    plt.xlabel("Frequency")
    plt.ylabel("Amplitude")
    plt.grid(True)
    plt.tight_layout()
    plt.show()

plot_frequency_spectrum(pred1, "Transformer w/o Sentiment")
plot_frequency_spectrum(pred2, "Transformer w/ Sentiment")
```

```
plot_frequency_spectrum(true_adj_close_wo, "True Closing Price")
```

